

# A Survey of High-Level Modeling and Simulation Methods for Modern Machine Learning Workloads

MICRO 2026 Submission – Confidential Draft – Do NOT Distribute!!

Anonymous Author(s)

Under Review

Anonymous

## Abstract

As machine learning workloads grow in scale and complexity—spanning training and inference for CNNs, transformers, mixture-of-experts models, and LLMs—architects and system designers need fast, accurate methods to predict their performance across diverse hardware platforms. This survey provides a comprehensive analysis of the tools and methods available for modeling and simulating the performance of ML workloads, covering analytical models, cycle-accurate simulators, trace-driven approaches, and ML-augmented hybrid techniques. We survey over 30 tools drawn from 53 papers across architecture venues (MICRO, ISCA, HPCA, ASPLOS) and systems venues (MLSys, OSDI, NSDI) published between 2016–2026, spanning DNN accelerator modeling (Timeloop, MAESTRO, Sparseloop), GPU simulation (GPGPU-Sim, Accel-Sim, NeuSight), distributed training simulation (ASTRA-sim, Lumos, SimAI), and LLM inference serving (VIDUR, Frontier, AMALI). We organize the literature along three dimensions—methodology type (analytical, simulation, ML-augmented, hybrid), target platform (accelerators, GPUs, distributed systems, edge devices), and abstraction level (kernel, model, system)—while additionally characterizing tools by workload coverage, revealing a pervasive CNN-validation bias. Our analysis reveals that hybrid approaches combining analytical structure with learned components achieve the best accuracy-speed trade-offs, while pure analytical models offer superior interpretability for design space exploration. We conduct hands-on reproducibility evaluations of five representative tools, finding that reproducibility varies dramatically: Docker-first tools score 8.5+/10 on our rubric while tools relying on serialized ML models risk becoming unusable. We identify key open challenges including cross-workload generalization beyond CNNs, composition of kernel-level predictions to end-to-end accuracy, and support for emerging architectures. This survey provides practitioners guidance for selecting appropriate modeling tools and researchers a roadmap for advancing the field of ML workload performance prediction.

## Keywords

ML workload performance prediction, DNN accelerator modeling, GPU simulation, distributed training simulation, LLM inference serving, design space exploration, survey

## 1 Introduction

Machine learning workloads—spanning training and inference for CNNs, transformers, mixture-of-experts models, and graph neural networks—have become the dominant consumers of compute

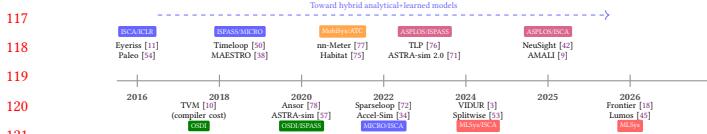
across datacenters and edge devices. The shift toward domain-specific architectures [22], from Google’s TPU [30, 31] to custom training accelerators, has created a heterogeneous hardware landscape where architects and system designers need fast, accurate performance predictions to navigate vast design spaces, select parallelization strategies, provision serving infrastructure, and optimize hardware-software co-design. Yet ML workloads pose unique modeling challenges: they exhibit diverse computational patterns (dense matrix operations in attention layers, sparse accesses in GNNs, communication-bound collective operations in distributed training) across this increasingly heterogeneous landscape of GPUs, TPUs, custom accelerators, and multi-device clusters.

A rich ecosystem of modeling and simulation tools has emerged to address these challenges, spanning a methodological spectrum from analytical models to cycle-accurate simulators to ML-augmented hybrid approaches. Analytical frameworks like Timeloop [50] and MAESTRO [38] model DNN accelerator performance through closed-form data movement analysis, achieving 5–10% error versus RTL at microsecond evaluation speed. Cycle-accurate simulators like GPGPU-Sim [4] and Accel-Sim [34] provide detailed GPU modeling but require hours per workload. Trace-driven simulators like ASTRA-sim [71] and VIDUR [3] target distributed training and LLM serving at system scale. ML-augmented approaches like NeuSight [42] learn performance functions from profiling data, achieving 2.3% error on GPU kernel prediction. Each methodology occupies a distinct point in the accuracy-speed-generality trade-off space.

Despite this rich tool landscape, no comprehensive survey organizes these methods from the perspective of the ML workload practitioner—the architect or engineer who needs to select a modeling tool for a specific design or deployment task. Existing surveys focus on ML *techniques* for performance modeling [65] or on specific hardware targets [50], leaving practitioners without guidance on which tools suit their needs across the full modeling spectrum. This survey fills that gap by providing a methodology-centric view of the tools and methods available for predicting ML workload performance.

We make the following contributions:

- A **methodology-centric taxonomy** organizing tools along three dimensions: methodology type (analytical, simulation, ML-augmented, hybrid), target platform (DNN accelerators, GPUs, distributed systems, edge devices), and abstraction level (kernel, model, system), with a quantitative coverage matrix identifying research gaps and a workload coverage analysis exposing the CNN-validation bias in the literature.
- A **systematic survey** of over 30 modeling tools drawn from 53 papers across architecture venues (MICRO, ISCA,



**Figure 1: Evolution of performance modeling tools for ML workloads (2016–2026).** Early analytical frameworks (EyeRISS, Paleo) gave way to systematic accelerator modeling (Timeloop, MAESTRO) and distributed training simulation (ASTRA-sim). ML-augmented approaches (TVM, Habitat, NeuSight) learn performance functions from data. Recent work targets LLM-specific modeling (VIDUR, AMALI, Frontier) and large-scale training prediction (Lumos).

HPCA, ASPLOS) and systems venues (MLSys, OSDI, NSDI) published between 2016–2026, using documented selection criteria.

- A **comparative analysis** examining trade-offs between accuracy, speed, generalization, and interpretability, with careful qualification of paper-reported accuracy claims and identification of cases where reported numbers are unverifiable.
- **Hands-on reproducibility evaluations** of representative tools with a 10-point rubric, and identification of **open challenges** including the CNN-to-transformer generalization gap, kernel-to-end-to-end error composition, and emerging accelerator support.

The remainder of this paper is organized as follows. Section 2 describes our survey methodology and positions this work relative to existing surveys. Section 3 provides background on ML workload characteristics and modeling fundamentals. Section 4 presents our classification taxonomy. Section 5 surveys approaches organized by target platform. Section 6 offers comparative analysis across key dimensions and a practitioner tool selection guide. Section 7 presents hands-on reproducibility evaluations. Section 8 discusses open challenges and future directions. Section 9 concludes.

Figure 1 illustrates the evolution of performance modeling tools for ML workloads, from early analytical frameworks through simulators to modern hybrid approaches.

## 2 Survey Methodology

We follow a systematic methodology for identifying, selecting, and classifying papers in this survey.

**Search strategy.** We searched ACM Digital Library, IEEE Xplore, Semantic Scholar, and arXiv using terms including “performance modeling DNN,” “DNN accelerator simulator,” “LLM inference prediction,” “distributed training simulation,” “neural network latency estimation,” and “ML workload performance.” We additionally performed backward/forward citation tracking from seminal works (Timeloop, ASTRA-sim, NeuSight) and monitored proceedings of target venues.

**Target venues.** Architecture: MICRO, ISCA, HPCA, ASPLOS. Systems: MLSys, OSDI, SOSP, NSDI. Related: NeurIPS, ICML, MobiSys, DAC, ISPASS.

**Inclusion criteria.** Papers must (1) propose or evaluate a tool or method for predicting performance of ML workloads (training or inference), (2) target at least one hardware platform (GPU, accelerator, distributed system, or edge device), and (3) include quantitative evaluation of prediction accuracy or modeling fidelity.

**Exclusion criteria.** We exclude (1) papers using ML for non-performance tasks (e.g., power estimation without latency), (2) papers modeling general-purpose (non-ML) workloads exclusively, and (3) papers without quantitative evaluation.

**Selection process.** Our initial search yielded 287 candidate papers. After title/abstract screening against inclusion criteria, 118 remained. Full-text review reduced the set to 53 papers that met all criteria. We additionally include 12 foundational works (gem5, roofline model, DRAMSim, etc.) as context for understanding the modeling landscape.

**Time period.** We cover papers published between 2016–2026, with foundational works from earlier years included for context.

**Classification.** We classify each paper along three dimensions: *methodology type* (analytical, cycle-accurate simulation, trace-driven simulation, ML-augmented, or hybrid), *target platform* (DNN accelerator, GPU, distributed system, edge device, or CPU), and *abstraction level* (kernel/operator, model/end-to-end, or system). We additionally characterize each tool by workload coverage, prediction targets, and reported accuracy metrics.

## 2.1 Related Surveys

Several surveys address adjacent topics. In the ML-for-systems space, Rakhshanfar and Zarandi [56] survey ML techniques for processor design space exploration, focusing on surrogate model construction rather than the tools available to ML practitioners. Sze et al. [66] provide a comprehensive treatment of DNN hardware architectures and dataflow optimization, establishing the conceptual framework on which analytical tools like Timeloop and MAESTRO are built; however, their scope is DNN accelerator design rather than cross-platform performance prediction. In GPU simulation, the gem5-gpu [6] and GPGPU-Sim [4] ecosystems have generated extensive evaluation literature, but no survey organizes GPU, accelerator, distributed, and edge modeling tools within a unified taxonomy. The MLPerf benchmark suites [47, 59] standardize ML workload measurement across hardware but focus on *measurement* rather than *prediction*—they provide ground truth data that performance models should target but do not survey the modeling tools themselves. Hennessy and Patterson [22] frame the current era as a “new golden age” for domain-specific architectures, motivating the need for performance prediction tools that span the heterogeneous hardware landscape, but do not survey these tools.

This survey differs from prior work in three ways: (1) it spans the full methodology spectrum from analytical to ML-augmented, rather than focusing on a single approach; (2) it covers all major target platforms (accelerators, GPUs, distributed systems, edge devices) rather than a single hardware class; and (3) it includes hands-on reproducibility evaluations that go beyond paper-reported accuracy claims. The closest prior work is the latency predictor study by Dudziak et al. [15], which systematically compares edge device predictors for NAS; we broaden the scope to the full platform and methodology landscape.

### 233 3 Background

234 This section provides background on the characteristics of ML work-  
 235 loads that make performance modeling challenging, and reviews  
 236 the fundamental approaches used to model them.

#### 238 3.1 ML Workload Characteristics

239 ML workloads present unique performance modeling challenges  
 240 compared to general-purpose programs. Modern ML frameworks  
 241 like PyTorch [52] and TensorFlow [1] define workloads as compu-  
 242 tation graphs of operators, providing a structured representation  
 243 that performance models can exploit.

244 **Computational structure.** ML workloads are composed of well-  
 245 defined operators (convolutions, matrix multiplications, attention  
 246 layers, normalization) with statically known shapes and data types.  
 247 This regularity enables analytical modeling of compute and data  
 248 movement, unlike branch-heavy general-purpose code. However,  
 249 modern architectures like mixture-of-experts (MoE) and dynamic  
 250 inference introduce input-dependent control flow that complicates  
 251 static analysis.

252 **Memory hierarchy sensitivity.** DNN accelerators employ spe-  
 253 cialized memory hierarchies with explicit data orchestration. The  
 254 mapping of tensor operations to hardware (dataflow, tiling, loop  
 255 ordering) critically determines performance. For LLM inference,  
 256 KV cache management dominates memory behavior, with cache  
 257 sizes scaling linearly with sequence length and batch size [39].

258 **Scale and distribution.** Large model training distributes com-  
 259 putation across thousands of GPUs using data, tensor, pipeline, and  
 260 expert parallelism [13]. Performance depends on the interplay be-  
 261 tween compute, memory bandwidth, and network communication—  
 262 requiring system-level modeling beyond single-device prediction.

263 **Distinct inference phases.** LLM inference exhibits qualita-  
 264 tively different phases: prefill (compute-bound, processing the full  
 265 prompt) and decode (memory-bound, generating tokens autore-  
 266 gressively) [53]. Effective modeling must capture both phases and  
 267 their interaction under batched serving [2, 74].

#### 270 3.2 Modeling Methodologies

271 We classify modeling approaches into four categories that form the  
 272 primary axis of our taxonomy.

273 **Analytical models** express performance as closed-form func-  
 274 tions of workload and hardware parameters. The roofline model [70]  
 275 bounds throughput by  $P = \min(\pi, \beta \cdot I)$ , where  $\pi$  is peak compute,  
 276  $\beta$  is memory bandwidth, and  $I$  is operational intensity. For DNN  
 277 accelerators, Timeloop [50] analytically computes data movement  
 278 costs across memory hierarchies for any valid mapping. Analytical  
 279 models provide microsecond evaluation and full interpretability,  
 280 but require manual derivation per architecture and struggle with  
 281 dynamic microarchitectural effects.

282 **Cycle-accurate simulators** model hardware at the register-  
 283 transfer level. gem5 [6] (CPUs), GPGPU-Sim [4] (GPUs), and Accel-  
 284 Sim [34] (modern GPUs) achieve detailed accuracy but suffer 1000–  
 285 10000× slowdown, making them impractical for full ML workload  
 286 evaluation. Sampling techniques (SimPoint [61], SMARTS [73]) re-  
 287 duce simulation time but were designed for general-purpose work-  
 288 loads and may not capture ML-specific patterns.

289 **Trace-driven simulation** uses execution traces as input rather  
 290 than full binary execution, enabling faster evaluation. ASTRA-  
 291 sim [71] models distributed training using Chakra execution traces [63]  
 292 with pluggable compute, memory, and network backends. VIDUR [3]  
 293 provides discrete-event simulation for LLM serving using kernel-  
 294 level profiles. This approach trades some fidelity for orders-of-  
 295 magnitude speedup over cycle-accurate simulation.

296 **ML-augmented approaches** learn performance functions from  
 297 profiling data. These range from simple models (random forests in  
 298 nn-Meter [77], XGBoost in TVM [10]) to deep learning (NeuSight [42])  
 299 and meta-learning (HELP [41]). ML-augmented approaches can cap-  
 300 ture complex non-linear relationships that elude analytical treat-  
 301 ment, but require training data and may not generalize beyond  
 302 their training distribution.

### 304 3.3 Problem Formulation

305 Performance modeling maps workload  $\mathcal{W}$  and hardware  $\mathcal{H}$  to a per-  
 306 formance metric  $y: \hat{y} = f(\mathcal{W}, \mathcal{H}; \theta)$ . Workloads are represented at  
 307 operator level (layer parameters), graph level (computation graphs),  
 308 IR level (compiler representations), or trace level (recorded runtime  
 309 behavior). Hardware is characterized by specifications, performance  
 310 counters, or learned embeddings.

311 **Prediction targets** include latency (execution time), throughput  
 312 (samples/second), energy (Joules per inference), and memory foot-  
 313 print. Multi-objective formulations enable Pareto-optimal design  
 314 selection.

315 **Accuracy metrics** vary across the literature: MAPE (scale-invariant  
 316 relative error), RMSE (penalizes large deviations), and rank corre-  
 317 lation (Kendall’s  $\tau$ ) for design space ordering. Direct comparison  
 318 across papers is limited by differences in benchmarks, hardware  
 319 targets, and evaluation protocols—a challenge we discuss in Sec-  
 320 tion 6.

## 4 Taxonomy

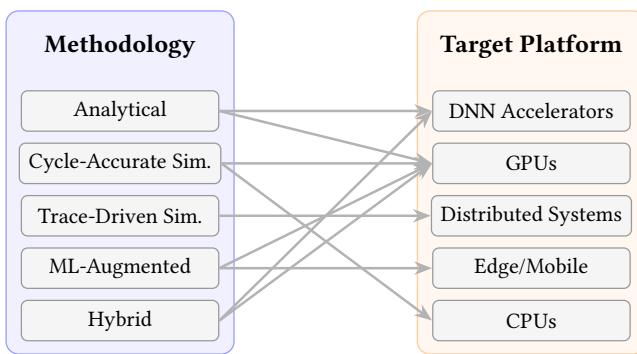
321 We organize the literature along three dimensions. The *primary axis*  
 322 is methodology type—how a tool predicts performance—because  
 323 methodology determines the fundamental trade-offs between accu-  
 324 racy, speed, interpretability, and data requirements. The *secondary  
 325 axes* are target platform and abstraction level, which together de-  
 326 termine the scope and applicability of each tool. We additionally  
 327 characterize tools by workload coverage, exposing a pervasive CNN-  
 328 validation bias in the literature.

329 Figure 2 illustrates the primary and secondary dimensions. Ta-  
 330 ble 1 provides a unified view combining the coverage matrix (num-  
 331 ber of surveyed tools per methodology–platform cell) with trade-off  
 332 profiles (evaluation speed, data requirements, interpretability, and  
 333 failure modes), with empty cells highlighting research gaps.

334 Table 1 reveals three structural observations. First, trace-driven  
 335 simulation is exclusively used for distributed systems—no surveyed  
 336 tool applies trace-driven methods to single-device GPU or accel-  
 337 erator modeling, despite the potential for trace-driven approaches to  
 338 avoid the slowdown of cycle-accurate simulation while retaining  
 339 more fidelity than analytical models. Second, edge/mobile devices  
 340 are served exclusively by ML-augmented approaches; the absence  
 341 of analytical or hybrid models for edge devices reflects the hard-  
 342 ware diversity problem but also represents a research gap, since  
 343

**Table 1: Methodology taxonomy: coverage matrix and trade-off profile.** Platform columns show the number of surveyed tools per cell; 0 indicates an explicit research gap. Speed, data requirements, and interpretability determine practical applicability; the failure mode column identifies the primary condition under which each methodology breaks down.

Methodology	DNN Accel.	GPU	Distrib. Systems	Edge/ Mobile	CPU	Eval. Speed	Data Req.	Interp.	Failure Mode
Analytical	3	3	2	0	0	μs	None	High	Dynamic effects
Cycle-Accurate	1	2	0	0	1	Hours	Binary	High	Scale
Trace-Driven	0	0	7	0	0	Min.	Traces	Med.	Trace fidelity
ML-Augmented	0	3	0	3	1	ms	Profiling	Low	Distrib. shift
Hybrid	1	3	0	0	1	ms	Mixed	Med.	Training domain



**Figure 2: Primary dimensions of the taxonomy for ML workload performance modeling.** The primary axis is methodology type (how performance is predicted); the secondary axis is target platform. The third dimension, abstraction level, is shown separately in Figure 3. Arrows show dominant pairings: analytical models for accelerators, cycle-accurate simulation for GPUs/CPUs, trace-driven simulation for distributed systems, and ML-augmented approaches for edge devices and compiler cost models.

hybrid approaches could combine the interpretability of analytical models with the adaptability of learned components. Third, no ML-augmented or hybrid tool specifically targets distributed system modeling—tools like VIDUR use ML internally for kernel prediction but are architecturally trace-driven simulators. The trade-off columns further show that methodologies cluster into two speed regimes: sub-millisecond (analytical, ML-augmented, hybrid) suitable for design space exploration, and minutes-to-hours (simulation, trace-driven) suitable for detailed validation.

## 4.1 Primary Axis: Methodology Type

The choice of methodology determines fundamental trade-offs between accuracy, evaluation speed, data requirements, and interpretability, as summarized in the right columns of Table 1.

**4.1.1 Analytical Models.** Analytical models express performance as closed-form functions of workload and hardware parameters. For DNN accelerators, Timeloop [50] models data movement across memory hierarchies for any valid loop-nest mapping, achieving 5–10% error versus RTL at 2000× speedup. MAESTRO [38] provides data-centric dataflow analysis using intuitive directives. Sparseloop [72]

extends to sparse tensor operations. Paleo [54] pioneered layer-wise analytical modeling for DNNs, decomposing networks into compute and communication components for distributed training prediction. AMALI [9] targets LLM inference on GPUs through improved memory hierarchy modeling.

Analytical models provide microsecond evaluation, full interpretability, and “what-if” design analysis. Their limitation is that they require manual derivation per architecture and may miss complex dynamic effects (e.g., memory contention, scheduling variability). AMALI’s 23.6% MAPE illustrates the accuracy ceiling of analytical approaches for complex GPU workloads—the residual error stems from dynamic microarchitectural effects that resist closed-form treatment (see §5.2 for detailed analysis).

**4.1.2 Cycle-Accurate Simulation.** Cycle-accurate simulators model hardware at register-transfer level, providing the highest fidelity. gem5 [6] (CPUs), GPGPU-Sim [4] (GPUs), and Accel-Sim [34] (modern NVIDIA GPUs, SASS-level trace-driven) achieve 0.90–0.97 IPC correlation. PyTorchSim [35] integrates PyTorch 2 with NPU simulation supporting custom RISC-V ISA and systolic arrays.

The primary limitation is speed: simulating a single ResNet-50 inference may require hours, making these tools impractical for design space exploration of ML workloads. Simulation sampling techniques (SimPoint [61], SMARTS [73], LoopPoint [60]) accelerate general-purpose workload simulation but are not specifically validated for ML workload patterns. Recent work on dissecting modern GPU cores [27] has improved Accel-Sim’s accuracy to 13.98% MAPE by reverse-engineering undocumented microarchitectural details. Note that the 0.90–0.97 IPC *correlation* metric can coexist with 20%+ absolute latency error for workloads with atypical occupancy patterns—correlation captures relative ordering fidelity but not absolute prediction accuracy.

**4.1.3 Trace-Driven Simulation.** Trace-driven approaches use recorded execution traces rather than full binary execution, enabling system-level modeling at practical speeds. ASTRA-sim [71] models distributed training end-to-end using Chakra execution traces [63], with pluggable compute, memory, and network backends, achieving 5–15% error versus real clusters. Echo [7] simulates distributed training at scale using analytical compute models with network simulation. Lumos [45] targets LLM training performance through trace-driven modeling, achieving 3.3% error on H100 GPUs.

For LLM inference serving, VIDUR [3] provides discrete-event simulation capturing prefill/decode phases, KV cache management, and request scheduling (Orca [74], Sarathi [2] strategies) with <5%

error. Frontier [18] extends to MoE and disaggregated inference with stage-centric simulation. SimAI [68] provides full-stack LLM training simulation achieving 98.1% alignment with production results at Alibaba Cloud scale.

These tools occupy a practical middle ground: fast enough for design exploration, detailed enough to capture system-level interactions that analytical models miss. Note that some tools in this category use ML internally (e.g., VIDUR uses random forests for kernel latency prediction), blurring the boundary with hybrid approaches—we classify by architectural design intent rather than implementation detail.

**4.1.4 ML-Augmented Models.** ML-augmented approaches learn performance functions entirely from profiling data, without embedding analytical domain knowledge. nn-Meter [77] uses random forest ensembles with kernel-level feature engineering for edge device latency prediction. LitePred [16] scales to 85 edge platforms using VAE-based intelligent sampling and transfer learning. HELP [41] formulates cross-hardware prediction as meta-learning, achieving adaptation with just 10 samples on new devices. TVM [10] and Anstor [78] use XGBoost/MLP cost models to guide compiler autotuning, with the TenSet dataset [79] (52M records) enabling pre-trained models.

ML-augmented approaches excel when sufficient profiling data is available and the training distribution matches deployment conditions. Their critical failure mode is *silent distribution shift*: a model trained on CNN kernels may produce confident but wrong predictions for transformer attention kernels, with no built-in mechanism to flag out-of-distribution inputs. nn-Meter’s paper-reported <1% MAPE cannot be independently verified, as the tool’s pre-trained predictors fail with current scikit-learn versions due to pickle serialization changes—a cautionary example of how ML-augmented approaches can become irreproducible and how unverifiable accuracy claims should be discounted by practitioners.

**4.1.5 Hybrid Analytical+ML Models.** Hybrid approaches combine analytical structure with learned components, achieving both interpretability and high accuracy. The analytical component provides a physics-based prior; the ML component learns residual corrections.

NeuSight [42] uses tile-based prediction mirroring CUDA’s execution model, achieving 2.3% error on GPT-3 inference. Concorde [49] fuses analytical models with learned corrections for CPU performance at 2% CPI error. Habitat [75] decomposes execution into analytically-modeled compute and memory components. Arch-Gym [37] connects ML optimization to analytical simulators for design space exploration.

The latency predictor study [15] demonstrates that hybrid approaches with transfer learning achieve 22.5% average improvement over baselines. Note that cross-tool accuracy comparisons require careful contextualization—we discuss methodological caveats (surrogate fidelity vs. real hardware error, evaluation-era fairness) in §5.2 and §5.5.

## 4.2 Secondary Axis: Target Platform

The target platform determines what performance effects must be modeled and constrains which methodologies are applicable.

**DNN Accelerators** (systolic arrays, dataflow architectures), from Google’s TPU [30, 31] to custom ASICs, are best served by analytical models (Timeloop, MAESTRO, Sparseloop) due to their regular, statically analyzable memory hierarchies and explicit dataflow control.

**GPUs** span the full methodology spectrum, from cycle-accurate (GPGPU-Sim, Accel-Sim) through analytical (AMALI, roofline [29, 70]) to hybrid (NeuSight, Habitat), reflecting the complexity of SIMD execution, warp scheduling, and memory coalescing.

**Distributed systems** are primarily served by trace-driven simulation (ASTRA-sim, VIDUR, Lumos, SimAI, Frontier) because system-level interactions (collective communication, pipeline parallelism, scheduling) cannot be captured by single-device models.

**Edge/mobile devices** are dominated by ML-augmented approaches (nn-Meter, LitePred, HELP) because the diversity of edge hardware makes per-device analytical modeling impractical.

**CPUs** for ML workloads are less studied because most ML training and inference runs on GPUs/accelerators. Concorde and GRANITE [65] target CPU performance but focus on general-purpose workloads rather than ML-specific patterns.

## 4.3 Secondary Axis: Abstraction Level

The abstraction level at which a tool operates determines what it can predict and where composition errors arise.

**Kernel/Operator-level** tools predict the latency of individual kernels or DNN operators (NeuSight, nn-Meter, TVM, GRANITE). They achieve the highest accuracy because the prediction scope is narrowly defined, but composing kernel predictions into end-to-end model latency introduces errors from memory allocation, kernel launch overhead, and inter-operator data movement.

**Model/End-to-End** tools predict full model inference or training time (Paleo, Habitat, AMALI, Lumos). They must account for graph-level effects (operator fusion, memory planning, scheduling) that kernel-level tools ignore, typically at the cost of higher error.

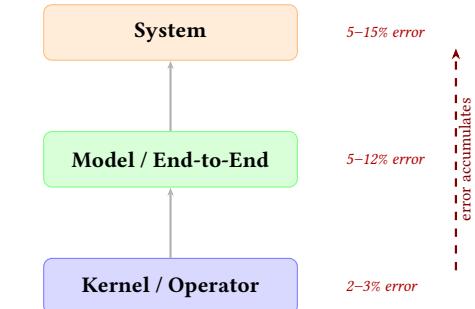
**System-level** tools predict multi-device or serving system performance (ASTRA-sim, VIDUR, SimAI, Frontier). They capture communication, scheduling, and resource contention effects but depend on the accuracy of their compute sub-models—creating a composition chain where kernel-level errors propagate through model-level to system-level predictions.

This three-level hierarchy makes explicit the *composition problem*: most tools operate at one level, but practitioners need predictions that span levels. The gap between kernel-level error (2–3% for NeuSight) and system-level error (5–15% for ASTRA-sim) reflects both the inherent difficulty of system modeling and the error accumulated through composition. Figure 3 illustrates this hierarchy and the error ranges at each level.

## 4.4 Workload Coverage

Table 2 characterizes the workload types on which each tool has been validated, exposing a pervasive CNN-validation bias.

The workload coverage table reveals that **no surveyed tool has been validated on diffusion models or dynamic inference workloads** (e.g., AI agents with tool use [36]). Only Frontier [18] has validated MoE support. For transformers, NeuSight, AMALI, VIDUR, and Frontier provide validated coverage, but each targets



**Figure 3: Abstraction level hierarchy and the composition problem.** Tools operate at one of three levels; composing predictions across levels accumulates error. Error ranges are representative values from surveyed papers.

**Table 2: Workload validation coverage.** ✓ = validated in the original paper; ○ = partial or indirect validation; — = no validation. Nearly all tools report accuracy on CNN workloads; transformer and MoE coverage is sparse. Empty columns (diffusion, dynamic inference) represent workload types with no validated performance modeling tools.

Tool	CNN	Trans- former	LLM Train	MoE	Diff.
Timeloop	✓	○	—	—	—
MAESTRO	✓	—	—	—	—
NeuSight	✓	✓	—	—	—
Habitat	✓	—	—	—	—
AMALI	—	✓	—	—	—
ASTRA-sim	✓	○	✓	—	—
VIDUR	—	✓	—	—	—
SimAI	—	—	✓	—	—
Lumos	—	—	✓	—	—
Frontier	—	✓	—	✓	—
nn-Meter	✓	—	—	—	—
LitePred	✓	—	—	—	—
HELP	✓	—	—	—	—
TVM/Ansor	✓	○	—	—	—

a different platform and abstraction level—no single tool offers validated transformer performance prediction across the full stack from kernel to system level. This workload coverage gap is the most actionable finding of our taxonomy: practitioners working with non-CNN workloads must either (a) accept unvalidated predictions from CNN-trained tools, (b) collect their own validation data, or (c) fall back to measurement.

## 5 Survey of Approaches

This section surveys performance modeling tools for ML workloads, organized by target platform. For each platform, we examine the modeling challenges, describe the available tools across methodology types, and critically analyze their strengths and limitations. Table 3 provides a comprehensive comparison.

### 5.1 DNN Accelerator Modeling

DNN accelerators employ specialized dataflows and memory hierarchies optimized for tensor operations [66]. The regularity of DNN computations makes this domain particularly amenable to analytical modeling.

**Analytical frameworks** dominate accelerator modeling. Timeloop [50] analytically computes data reuse, latency, and energy from loopnest representations, achieving 5–10% error versus RTL simulation at 2000× speedup. It provides reference outputs for standard accelerator designs (Eyeriss [11], Simba) with deterministic results—a key reproducibility strength. MAESTRO [38] offers data-centric dataflow directives that simplify specification but is less precise than Timeloop for detailed energy modeling. Sparseloop [72] extends Timeloop to sparse tensor operations by modeling the interaction between sparsity patterns (structured vs. unstructured), compression formats (CSR, bitmap), and hardware decompression/intersection units. This is critical for efficient transformer inference where attention matrices exhibit structured sparsity, but Sparseloop assumes static, known sparsity distributions—dynamic sparsity patterns from techniques like token pruning or dynamic routing in MoE models fall outside its modeling capability.

**Simulation approaches.** PyTorchSim [35] integrates PyTorch 2 with cycle-accurate NPU simulation, supporting custom RISC-V ISA and systolic arrays with configurable memory hierarchies. Unlike standalone accelerator simulators, PyTorchSim directly consumes PyTorch computation graphs, eliminating the manual workload translation step that introduces errors and limits adoption. However, it does not report accuracy against real hardware, and its cycle-accurate approach inherits the speed limitations of simulation-based methods, making it impractical for large-model evaluation.

**ML-augmented design.** ArchGym [37] connects ML optimization algorithms to analytical simulators for design space exploration. Its reported 0.61% RMSE measures how faithfully the ML surrogate reproduces the simulator’s predictions—not accuracy against real hardware. This distinction matters: surrogate fidelity enables fast DSE but does not validate the underlying simulator’s accuracy.

**Emerging accelerators.** Processing-in-memory (PIM) architectures present fundamentally different modeling challenges, as they blur the compute-memory boundary that conventional frameworks assume. Early PIM modeling tools [23, 28, 40, 51] target attention acceleration and heterogeneous PIM-GPU co-simulation, but none report accuracy against real PIM hardware—we discuss PIM modeling gaps further in Section 8.3.

**Synthesis.** Accelerator modeling is the most mature subdomain surveyed, with Timeloop’s analytical framework achieving a favorable balance of accuracy (5–10% error), speed, and interpretability that has made it the de facto standard for accelerator design space exploration. The progression from Timeloop through Sparseloop to PIM-aware tools illustrates a recurring pattern: each extension addresses a new workload characteristic (sparsity, near-memory compute) but adds modeling complexity that erodes the simplicity advantage of analytical approaches. The key gap is cycle-accurate validation—ArchGym and PyTorchSim provide simulation-based alternatives, but neither validates against manufactured silicon, leaving the accuracy of all accelerator modeling tools ultimately anchored to RTL comparisons rather than measured hardware.

**Table 3: Summary of surveyed performance modeling tools for ML workloads, organized by target platform. Methodology:** A=Analytical, S=Simulation, T=Trace-driven, M=ML-augmented, H=Hybrid. \*Accuracy measures surrogate-vs-simulator fidelity, not real hardware error. †Reported accuracy unverifiable due to reproducibility issues. ‡No accuracy baseline against real hardware reported.

Tool	Platform	Method	Target	Accuracy	Speed	Key Capability
<i>DNN Accelerator Modeling</i>						
Timeloop [50]	NPU	A	Latency/Energy	5–10%	μs	Loop-nest DSE
MAESTRO [38]	NPU	A	Latency/Energy	5–15%	μs	Data-centric directives
Sparseloop [72]	NPU	A	Sparse tensors	5–10%	μs	Compression modeling
PyTorchSim [35]	NPU	S	Cycle-accurate	N/A <sup>‡</sup>	Hours	PyTorch 2 integration
ArchGym [37]	Multi	H	Multi-objective	0.61%*	ms	ML-aided DSE
<i>GPU Performance Modeling</i>						
Accel-Sim [34]	GPU	S	Cycle-accurate	10–20%	Hours	SASS trace-driven
GPGPU-Sim [4]	GPU	S	Cycle-accurate	10–20%	Hours	CUDA workloads
AMALI [9]	GPU	A	LLM inference	23.6%	ms	Memory hierarchy
NeuSight [42]	GPU	H	Kernel/E2E latency	2.3%	ms	Tile-based prediction
Habitat [75]	GPU	H	Training time	11.8%	Per-kernel	Wave scaling
<i>Distributed Training and LLM Serving</i>						
ASTRA-sim [71]	Distributed	T	Training time	5–15%	Minutes	Collective modeling
SimAI [68]	Distributed	T	Training time	1.9%	Minutes	Full-stack simulation
Lumos [45]	Distributed	T	LLM training	3.3%	Minutes	H100 training
VIDUR [3]	GPU cluster	T	LLM serving	<5%	Seconds	Prefill/decode phases
Frontier [18]	Distributed	T	MoE inference	—	Minutes	Stage-centric sim.
TrioSim [43]	Multi-GPU	T	DNN training	N/A <sup>‡</sup>	Minutes	Lightweight multi-GPU
<i>Edge Device Modeling</i>						
nn-Meter [77]	Edge	M	Latency	<1%†	ms	Kernel detection
LitePred [16]	Edge	M	Latency	0.7%	ms	85-platform transfer
HELP [41]	Multi	M	Latency	1.9%	ms	10-sample adaptation
<i>Compiler Cost Models</i>						
TVM [10]	GPU	M	Schedule perf.	~15%	ms	Autotuning guidance
Anstor [78]	GPU	M	Schedule perf.	~15%	ms	Program sampling
TLP [76]	GPU	M	Tensor program	<10%	ms	Transformer cost model

## 5.2 GPU Performance Modeling

GPUs dominate ML training and inference, making accurate GPU performance prediction critical. GPU modeling must account for SIMT execution, warp scheduling, memory coalescing, and workload-dependent occupancy effects.

**Cycle-accurate simulation.** GPGPU-Sim [4] and Accel-Sim [34] achieve 0.90–0.97 IPC correlation through detailed microarchitectural modeling. Recent work reverse-engineering modern GPU cores [27] has improved Accel-Sim to 13.98% MAPE by modeling previously undocumented features. However, 1000–10000× slowdown makes these tools impractical for full ML workloads at production scale.

**Analytical models.** The roofline model [70] provides a useful upper bound but misses occupancy and memory hierarchy effects. Roofline-LLM [29] extends roofline analysis to LLM inference. AMALI [9] reduces GPU LLM inference MAPE from 127% (prior analytical baselines) to 23.6% through improved memory hierarchy modeling. The residual 23.6% error reflects the fundamental difficulty of analytically modeling GPU dynamic behavior (warp scheduling, L2 cache contention, bank conflicts) rather than a quality limitation.

**Hybrid learned models.** NeuSight [42] introduces tile-based prediction that mirrors CUDA’s execution model, achieving 2.3% MAPE on GPT-3 inference across H100, A100, and V100 GPUs. Habitat [75] decomposes execution into analytically-modeled compute and memory components using wave scaling analysis, achieving 11.8% error for cross-GPU transfer (e.g., V100→A100). Its key insight is that compute and memory bandwidth scale independently across GPU generations, enabling prediction on new hardware from profiling data on existing hardware. However, Habitat requires source GPU profiling, limiting its use for pre-silicon design exploration, and its wave scaling model assumes that GPU occupancy patterns remain similar across generations—an assumption that breaks for workloads with qualitatively different memory access patterns (e.g., KV cache in LLM decode vs. activation reuse in CNN training). Direct comparison between NeuSight and Habitat requires caution: NeuSight evaluates on 2023–2025 hardware (H100) with LLM workloads, while Habitat was designed for earlier GPUs with CNN/RNN workloads—the reported “50× improvement” reflects different evaluation conditions rather than purely methodological advances.

**LLM-specific modeling.** LLM execution exhibits qualitatively different prefill (compute-bound) and decode (memory-bound) phases [53,

80]. VIDUR [3] provides discrete-event simulation for LLM serving  
 81 systems, capturing request scheduling strategies (Orca [74],  
 82 Sarathi [2]) with <5% error. LIFE [17] offers hardware-agnostic analytical  
 83 LLM inference modeling by decomposing inference into compute and memory-access components that can be parameterized  
 84 for arbitrary hardware specifications, enabling performance prediction without hardware-specific profiling. Its hardware-agnostic  
 85 design enables pre-silicon evaluation of new accelerators for LLM workloads, but the analytical approach shares the same accuracy  
 86 limitations as AMALI when applied to GPUs with complex dynamic behavior. HERMES [5] targets heterogeneous multi-stage LLM inference  
 87 pipelines where different model components (embedding, attention, FFN) execute on different hardware, modeling the inter-stage data transfer and load balancing that arise in disaggregated serving architectures. Emerging work uses LLMs for GPU kernel  
 88 performance prediction: Omniwise [21] achieves 90% of predictions within 10% error on AMD MI250/MI300X, and SwizzlePerf [67]  
 89 achieves 2.06× speedup through hardware-aware spatial optimization.  
 90

**Compiler cost models.** TVM [10] and Ansor [78] use ML cost models (XGBoost, MLP) to guide autotuning, achieving ~15% MAPE. The TenSet dataset [79] (52M records) enables pre-trained models that accelerate autotuning 10×. TLP [76] uses deep learning (transformer-based architecture) for tensor program cost modeling, specifically designed for the irregular computation patterns in transformer workloads that challenge traditional XGBoost cost models; it achieves <10% MAPE on transformer operators where TVM’s default cost model shows higher error, demonstrating that workload-specific cost models can significantly improve autotuning for non-CNN workloads. SynPerf [69] takes a complementary approach, using performance models to guide GPU kernel synthesis rather than merely evaluating existing kernels. These tools prioritize ranking accuracy for schedule selection over absolute error.

**Synthesis.** GPU modeling exhibits the widest methodological spread of any platform category: cycle-accurate simulation (Accel-Sim), analytical models (AMALI, roofline), hybrid learned approaches (NeuSight, Habitat), and LLM-based predictors (Omniwise) all target the same hardware with error rates spanning 2%–24%. This diversity reflects the fundamental tension in GPU modeling: the microarchitectural complexity that makes GPUs powerful also makes them hard to model analytically, while the rapid hardware evolution that motivates prediction also invalidates training data for learned approaches. NeuSight’s tile-based decomposition currently offers the best accuracy–speed trade-off for LLM workloads, but its reliance on per-GPU profiling limits pre-silicon use—a gap that analytical approaches like AMALI and LIFE fill despite higher error. The compiler cost model ecosystem (TVM, TLP, SynPerf) represents a distinct use case where relative ranking matters more than absolute prediction, suggesting that evaluation criteria should be workload-aware.

### 5.3 Distributed Training and LLM Serving

Distributed systems introduce communication overhead, synchronization barriers, and parallelism strategy choices. Modern large

model training uses multiple parallelism dimensions: tensor parallelism splits individual layers across GPUs [62], pipeline parallelism distributes layers across pipeline stages [26], and memory-efficient optimizers like ZeRO [55] partition optimizer states across data-parallel workers. Performance depends on the interplay between compute, memory, and network—requiring system-level modeling.

**Training simulation.** ASTRA-sim [71] provides end-to-end distributed training simulation using Chakra execution traces [63], with validated HGX-H100 configurations and pluggable network backends. It achieves 5–15% error versus real clusters and enables exploration of parallelization strategies at scale. SimAI [68] provides full-stack LLM training simulation at Alibaba Cloud scale, modeling compute, memory, network, and collective communication in an integrated framework that achieves 1.9% MAPE versus production training runs. Its key differentiator is validation against production training runs at datacenter scale—most simulators validate at 4–64 GPU configurations, whereas SimAI validates against thousands of GPUs where network congestion and load imbalance effects dominate. Lumos [45] targets LLM training through trace-driven modeling, achieving 3.3% error on H100 GPUs by capturing gradient accumulation, optimizer states, and activation checkpointing. Echo [7] combines analytical compute models with packet-level network simulation for collective communication evaluation, though it does not report end-to-end accuracy against real hardware. TrioSim [43] offers lightweight multi-GPU simulation through selective fidelity, enabling rapid what-if analysis for multi-GPU configurations but without real-hardware accuracy baselines. PRISM [19] produces prediction intervals rather than point estimates at 10K+ GPU scale, capturing the stochastic performance variation (network congestion, stragglers) that deterministic simulators miss.

**Scaling and parallelism.** The choice of parallelism strategy (data, tensor, pipeline, expert) critically impacts performance. Paleo [54] pioneered analytical estimation of training time by decomposing workloads into compute and communication components. MAD Max [25] decomposes training time into compute, communication, and memory components per parallelism dimension, enabling rapid analytical evaluation of parallelism configurations without simulation. The Llama 3 scaling study [13] documents 4D parallelism at 16K H100 GPUs, providing ground truth for simulator validation. Sailor [64] addresses automated parallelism selection over heterogeneous clusters, where GPUs of different generations or types must be jointly scheduled—a problem that most simulators cannot model because they assume homogeneous hardware.

**Inference serving.** VIDUR [3] simulates LLM inference serving with scheduling strategies (vLLM [39], Orca [74], Sarathi [2]) without requiring GPU hardware. Frontier [18] extends to MoE and disaggregated inference with stage-centric simulation. ThrottLL’eM [32] models the interaction between GPU power management (frequency throttling, power capping) and LLM inference performance, addressing a dimension that most tools ignore: real GPU deployments operate under power and thermal constraints that reduce effective performance below theoretical peaks. By modeling throttling effects, ThrottLL’eM enables energy-efficient inference scheduling that trades latency headroom for power savings—a critical concern for datacenter operators where energy costs dominate TCO. Recent LLM inference optimizations also change the

929 performance characteristics that simulators must capture: for example, MEDUSA [8] introduces speculative decoding that transforms  
 930 sequential token generation into parallel verification, fundamentally altering the compute-to-memory ratio that models like VIDUR  
 931 assume. Such optimizations illustrate a moving-target challenge:  
 932 performance models must track not just hardware evolution but  
 933 algorithmic innovations that restructure execution patterns. These  
 934 tools collectively enable infrastructure planning and scheduling  
 935 algorithm comparison at scale.

936 **Memory system interactions.** Memory increasingly dominates ML performance: KV cache management is the key LLM  
 937 serving bottleneck (vLLM’s PagedAttention [39] achieves 2–4×  
 938 throughput improvement), and VIDUR models cache allocation and  
 939 eviction at the system level. Low-level memory simulators (DRAM-  
 940 Sim3 [44], Ramulator 2 [46]) integrate with tools like Accel-Sim  
 941 rather than being used standalone for ML workloads.

942 **Synthesis.** Distributed system modeling is the fastest-growing  
 943 subdomain, with six new tools published since 2024 (SimAI, Lu-  
 944 mos, Echo, TrioSim, PRISM, Sailor). This surge reflects the  
 945 practical urgency: training runs on thousands of GPUs cost millions  
 946 of dollars, making pre-deployment performance prediction eco-  
 947 nomically critical. The tools bifurcate into two design philosophies:  
 948 *trace-driven fidelity* (ASTRA-sim, SimAI) replays recorded execu-  
 949 tion traces through pluggable backends for maximum realism, while  
 950 *analytical decomposition* (Paleo, MAD Max) trades fidelity for speed  
 951 and interpretability. PRISM’s probabilistic approach represents an  
 952 emerging third path, acknowledging that large-scale systems are  
 953 inherently stochastic. The inference serving tools (VIDUR, Frontier,  
 954 ThrottLL’eM) face a distinct challenge: algorithmic innovations  
 955 like speculative decoding [8] continuously alter the performance  
 956 characteristics that models must capture, creating a moving-target  
 957 problem absent in training simulation.

## 961 5.4 Edge Device Modeling

962 Edge devices impose strict power, memory, and latency constraints.  
 963 The diversity of edge hardware (mobile CPUs, GPUs, NPUs, DSPs)  
 964 makes per-device analytical modeling impractical, leading to ML-  
 965 augmented approaches.

966 nn-Meter [77] uses random forest ensembles with kernel-level  
 967 feature engineering, reporting <1% MAPE. However, this claim is  
 968 currently unverifiable: the tool’s pre-trained predictors fail with  
 969 modern scikit-learn versions due to pickle serialization changes,  
 970 scoring only 3/10 in our reproducibility evaluation. LitePred [16]  
 971 scales to 85 edge platforms using VAE-based intelligent sampling  
 972 and transfer learning, achieving 0.7% MAPE with under one hour of  
 973 adaptation per device. Its key innovation is intelligent sample selec-  
 974 tion: rather than profiling all operators on a new device, LitePred  
 975 uses a VAE to identify the most informative operators to profile,  
 976 reducing adaptation cost by an order of magnitude. However, the  
 977 “85 platforms” are predominantly ARM-based mobile CPUs and  
 978 GPUs—the diversity within this evaluation set is unclear, and trans-  
 979 fer to fundamentally different accelerators (NPUs, DSPs) likely  
 980 degrades significantly. HELP [41] formulates cross-hardware pre-  
 981 diction as meta-learning with MAML-style adaptation, achieving  
 982 1.9% MAPE with just 10 measurement samples on new devices. The  
 983 10-sample adaptation is compelling for rapid deployment but raises  
 984

985 a selection problem: which 10 operators to profile depends on the  
 986 target workload, and suboptimal sample selection can significantly  
 987 degrade accuracy on workload-critical operators not represented  
 988 in the adaptation set. ESM [48] provides a systematic framework  
 989 for building effective surrogate models for hardware-aware neural  
 990 architecture search (NAS), evaluating multiple ML architectures  
 991 (MLPs, gradient-boosted trees, GNNs) as latency surrogates across  
 992 different hardware platforms. Its key finding is that model archi-  
 993 tecture choice matters less than training data quality and feature  
 994 engineering—well-tuned random forests match or outperform deep  
 995 learning surrogates, challenging the assumption that more complex  
 996 models yield better hardware-aware NAS performance. This result  
 997 has practical implications for the ML-augmented tools surveyed  
 998 here: the accuracy gains from sophisticated model architectures  
 999 may be marginal compared to improvements in profiling data col-  
 1000 lection.

1001 The latency predictor study [15] provides the most systematic  
 1002 evaluation across approaches, showing transfer learning provides  
 1003 22.5% average improvement, up to 87.6% on challenging cross-  
 1004 platform transfers.

1005 **Synthesis.** Edge modeling stands apart from the other platform  
 1006 categories in being dominated by ML-augmented approaches—the  
 1007 hardware diversity makes analytical modeling impractical, so the  
 1008 field has converged on learning latency functions from profiling  
 1009 data. The central challenge is *generalization*: each tool (nn-Meter,  
 1010 LitePred, HELP) proposes a different strategy for adapting to new  
 1011 devices with minimal profiling, yet ESM’s finding that simple ran-  
 1012 dom forests match deep learning surrogates suggests the field may  
 1013 be over-investing in model complexity relative to data quality. The  
 1014 reproducibility crisis exemplified by nn-Meter (pre-trained models  
 1015 that become unusable across library versions) serves as a warning  
 1016 for the entire ML-augmented approach: accuracy claims are only  
 1017 valuable if the tools remain functional over time.

## 1019 5.5 Cross-Cutting Challenges

1020 Several challenges cut across platform categories.

1021 **CNN-to-transformer gap.** Nearly all reported accuracy num-  
 1022 bers are measured on CNN workloads. Performance on transform-  
 1023 ers, MoE, and diffusion models is less well characterized. NeuSight  
 1024 is a notable exception, evaluating on GPT-3 inference, but most  
 1025 tools lack validated transformer support.

1026 **Kernel-to-end-to-end composition.** Many tools predict kernel-  
 1027 level performance (nn-Meter, NeuSight), but composing kernel pre-  
 1028 dictions into accurate end-to-end estimates is an unsolved problem.  
 1029 Memory allocation, kernel launch overhead, and inter-operator  
 1030 data movement introduce errors that compound across layers.

1031 **Static vs. profiling-based approaches.** A fundamental practi-  
 1032 cal divide exists between tools that predict from static specifications  
 1033 only (Timeloop, MAESTRO, Paleo) and those requiring runtime  
 1034 profiling data (Habitat, nn-Meter, HELP). Static approaches enable  
 1035 pre-silicon evaluation and NAS; profiling-based approaches achieve  
 1036 higher accuracy on existing hardware. This distinction is often more  
 1037 practically relevant than the analytical-vs-ML divide.

1038 **Design patterns in successful tools.** Across all platform cate-  
 1039 gories, the most effective tools share common design choices. First,

1045 *structural decomposition* that mirrors hardware execution consistently outperforms black-box approaches: Timeloop’s loop-nest representation captures accelerator dataflows, NeuSight’s tile-based decomposition mirrors CUDA execution, and VIDUR’s prefill/decode phase separation reflects the actual memory-vs-compute regime shift in LLM serving. These tools succeed because their abstractions encode domain knowledge about *why* performance varies, not just correlations. Second, tools with the strongest practical adoption provide modular, pluggable backends—ASTRA-sim supports both analytical and ns-3 network simulation, and VIDUR integrates multiple scheduling algorithms (Orca [74], Sarathi [2])—allowing users to trade accuracy for speed depending on the evaluation scenario. Third, robust reproducibility correlates with sustained community use: Timeloop (9/10 reproducibility in our evaluation) and ASTRA-sim (8.5/10) have mature Docker support and deterministic outputs, while tools with higher reported accuracy but poor reproducibility (nn-Meter claims <1% MAPE but scored 3/10 due to dependency failures) see declining adoption. Our reproducibility findings suggest that *verifiable moderate accuracy* matters more to practitioners than *unverifiable high accuracy*.

1065 **Accuracy claims require careful contextualization.** Comparing accuracy numbers across the surveyed tools is misleading without accounting for problem difficulty. Predicting single-kernel latency on a known device (nn-Meter, LitePred) is fundamentally easier than predicting end-to-end distributed training time across thousands of GPUs (ASTRA-sim, SimAI [68]), yet the former reports sub-1% error while the latter reports 5–15%. Similarly, ArchGym’s 0.61% RMSE measures surrogate-vs-simulator fidelity—a regression task over a smooth design space—not prediction of real hardware behavior. Even within a single platform, methodology choice constrains achievable accuracy: AMALI’s 23.6% analytical error versus NeuSight’s 2.3% ML-based error on GPU LLM inference reflects the fundamental difficulty of capturing GPU dynamic behavior (warp scheduling, cache contention) in closed-form models, not a quality gap between tools. These comparisons highlight that *problem difficulty* and *what is being measured* must be specified alongside any accuracy number; Section 6 provides a structured analysis accounting for these factors.

1083 **The gap between model output and practitioner needs.** A recurring limitation across all platform categories is the mismatch between what tools predict and what deployment decisions require. Most tools predict *compute latency* or *throughput* for individual operations, but practitioners need *time-to-accuracy* for training (which depends on convergence, not just iteration time), *tail latency under load* for serving (which depends on scheduling and queuing effects), and *operational cost* for capacity planning (which depends on utilization, failures, and thermal throttling [32]). Only a few tools partially bridge this gap: VIDUR models scheduling-level effects, Lumos [45] captures training-specific overheads like gradient accumulation and activation checkpointing, and PRISM [19] produces prediction intervals rather than point estimates to reflect inherent system variability. Closing this gap—connecting component-level performance predictions to system-level deployment metrics—remains the most impactful direction for future tool development.

## 6 Comparison and Analysis

We analyze trade-offs across methodology types along four dimensions: accuracy, speed, generalization, and interpretability. Table 4 summarizes key characteristics.

### 6.1 Accuracy by Problem Difficulty

Rather than comparing accuracy numbers directly (which is misleading across different benchmarks, metrics, and hardware), we organize results by problem difficulty.

**Accelerator dataflow modeling** is the most amenable to accurate prediction because computations are regular and memory access patterns are statically determined. Timeloop achieves 5–10% error against RTL through purely analytical means.

**Single-GPU kernel prediction** for known architectures achieves 2–12% error through hybrid approaches (NeuSight, Habitat) that embed hardware-specific inductive biases.

**Distributed system-level prediction** achieves 2–15% error through trace-driven simulation (SimAI 1.9%, Lumos 3.3%, ASTRA-sim 5–15%), reflecting the challenge of modeling compute-communication interaction.

**Cross-platform edge prediction** achieves 0.7–2% error (LitePred, HELP) but requires per-device profiling data, trading generality for accuracy.

**GPU analytical modeling** remains the most difficult, with AMALI’s 23.6% representing the current state of the art for purely analytical GPU LLM inference prediction—a problem where dynamic microarchitectural effects resist closed-form treatment.

Figure 4 visualizes reported accuracy (MAPE) across tools, grouped by methodology type. The pattern is clear: hybrid approaches achieve the lowest error on GPU workloads, trace-driven simulators cluster at 2–15% for distributed systems, and analytical models trade accuracy for speed and interpretability. Note that direct comparison across tools is approximate because accuracy numbers are measured on different benchmarks, workloads, and hardware targets; the figure illustrates methodology-level trends rather than head-to-head rankings.

Figure 5 plots representative tools on two axes—evaluation speed versus reported accuracy—revealing the fundamental trade-off space. Analytical models (upper-left) achieve the fastest evaluation but sacrifice accuracy on complex workloads. Cycle-accurate simulators (lower-right) provide the highest fidelity but at impractical speeds. Hybrid approaches (NeuSight, Concorde) occupy the desirable upper-left region: fast evaluation *and* low error, though at the cost of training data requirements and reduced interpretability compared to analytical models. Trace-driven simulators span a wide range, from VIDUR’s seconds-scale LLM serving simulation to ASTRA-sim’s minutes-scale distributed training, reflecting the diversity of system-level modeling targets.

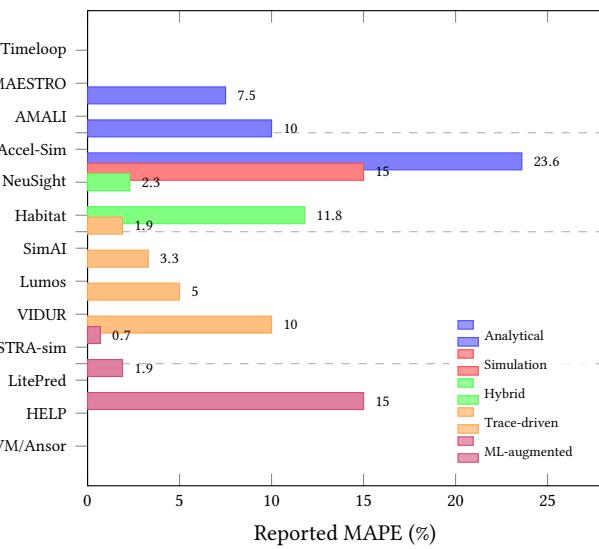
### 6.2 Generalization Challenges

**Workload generalization.** Nearly all reported accuracy numbers are measured on CNN workloads. Cross-workload-type transfer (CNN→transformer) remains largely unvalidated. NeuSight is a notable exception, evaluating on LLM workloads, but most edge device predictors (nn-Meter, LitePred, HELP) are validated primarily on CNNs.

1103  
1104  
1105  
1106  
1107  
1108  
1109  
1110  
1111  
1112  
1113  
1114  
1115  
1116  
1117  
1118  
1119  
1120  
1121  
1122  
1123  
1124  
1125  
1126  
1127  
1128  
1129  
1130  
1131  
1132  
1133  
1134  
1135  
1136  
1137  
1138  
1139  
1140  
1141  
1142  
1143  
1144  
1145  
1146  
1147  
1148  
1149  
1150  
1151  
1152  
1153  
1154  
1155  
1156  
1157  
1158  
1159  
1160

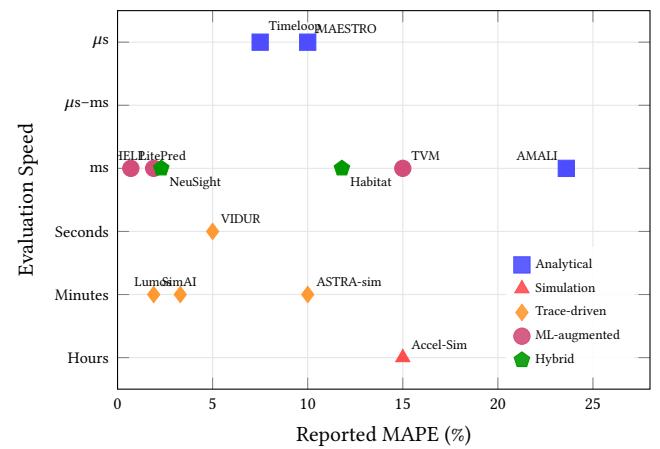
**Table 4: Comparative analysis of representative tools across key dimensions. Accuracy figures are as reported in original papers; direct comparison is limited by differences in benchmarks, workloads, hardware targets, and evaluation protocols.**<sup>†</sup>Unverifiable. <sup>\*</sup>Surrogate fidelity, not hardware accuracy.

Tool	Methodology	Accuracy (reported)	Setup Cost	Generalization	Interpretability	Eval. Speed
Timeloop [50]	Analytical	5–10%	Arch spec only	Any accelerator	High	$\mu\text{s}$
MAESTRO [38]	Analytical	5–15%	Arch spec only	Any accelerator	High	$\mu\text{s}$
AMALI [9]	Analytical	23.6% MAPE	None	GPU LLM inference	High	ms
Accel-Sim [34]	Simulation	10–20%	GPU binary	GPU-specific	High	Hours
ASTRA-sim [71]	Trace-driven	5–15%	Execution trace	Configurable	Medium	Minutes
VIDUR [3]	Trace-driven	<5%	Kernel profiles	LLM-specific	High	Seconds
SimAI [68]	Trace-driven	1.9%	Full-stack setup	LLM training	Medium	Minutes
Lumos [45]	Trace-driven	3.3%	Execution trace	LLM training	Medium	Minutes
nn-Meter [77]	ML-augmented	<1% <sup>†</sup>	1K samples/kernel	Device-specific	Medium	ms
LitePred [16]	ML-augmented	0.7% MAPE	100 samples/device	85+ devices	Low	ms
HELP [41]	ML-augmented	1.9% MAPE	10 samples/device	Cross-platform	Low	ms
TVM [10]	ML-augmented	~15% MAPE	10K+	Operator-level	Medium	ms
NeuSight [42]	Hybrid	2.3% MAPE	Pre-trained	Cross-GPU	Medium	ms
Habitat [75]	Hybrid	11.8% MAPE	Online profiling	Cross-GPU	Medium	Per-kernel
ArchGym [37]	Hybrid	0.61% RMSE*	Simulation runs	Arch-specific	Medium	ms
Concorde [49]	Hybrid	2% CPI	Training corpus	Cross- $\mu$ arch	Medium	ms

**Figure 4: Reported accuracy (MAPE) of surveyed tools, grouped by methodology type. Range midpoints are used where ranges are reported (e.g., 7.5% for Timeloop’s 5–10%). Cross-tool comparison is approximate due to differing benchmarks, workloads, and hardware targets.**

**Hardware generalization.** Three strategies show promise: meta-learning (HELP with 10-sample adaptation), feature-based transfer (LitePred across 85 devices), and analytical decomposition (Habitat separating compute/memory scaling). Cross-family transfer (GPU→TPU→PIM) remains unsolved.

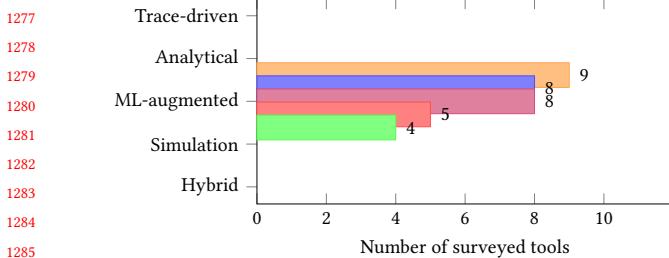
**Temporal generalization.** Software stack evolution (framework updates, driver changes, compiler optimizations) invalidates

**Figure 5: Speed–accuracy trade-off for surveyed tools. The y-axis represents evaluation speed (higher is faster); the x-axis shows reported MAPE (lower is better). The desirable region is the upper-left quadrant (fast and accurate). Hybrid approaches cluster in the fast-and-accurate region; analytical models are fastest but less accurate on complex workloads; cycle-accurate simulation is slowest but captures microarchitectural detail.**

trained models over time. No surveyed tool addresses continual learning for evolving software environments.

### 6.3 Interpretability and Design Insight

A key advantage of analytical models is actionable design insight. Timeloop identifies data movement bottlenecks; MAESTRO reveals suboptimal dataflow choices; VIDUR exposes scheduling inefficiencies. These insights directly guide design decisions.



**Figure 6: Distribution of surveyed tools by methodology type.** Trace-driven simulation dominates due to the recent growth of distributed training and LLM serving tools. Hybrid approaches are the least represented despite their strong accuracy results.

ML-augmented approaches (nn-Meter, HELP) provide feature importance rankings but limited causal understanding. Hybrid approaches (NeuSight, Concorde) offer partial interpretability through their analytical components. The interpretability gap is practically significant: architects need to understand *why* a design is slow, not just predict *that* it is slow.

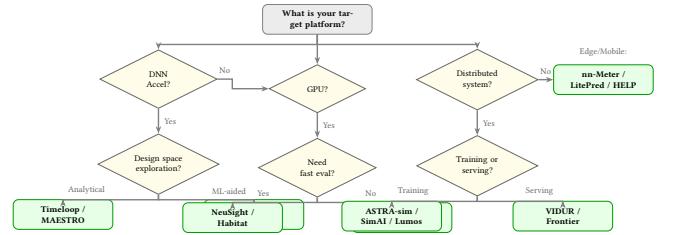
#### 6.4 Methodology Distribution

Figure 6 shows the distribution of surveyed tools across methodology types. Trace-driven simulation is the most common approach (9 tools), driven by the recent proliferation of distributed training and LLM serving simulators. Analytical and ML-augmented approaches are equally represented (8 tools each), reflecting the field’s split between interpretable physics-based models and data-driven prediction. Hybrid approaches remain the smallest category (4 tools), despite achieving the best accuracy—suggesting significant room for future work combining analytical structure with learned components.

#### 6.5 Practitioner Tool Selection Guide

To address the gap between taxonomy and actionable guidance, Figure 7 presents a decision flowchart for practitioners selecting a performance modeling tool. The flowchart captures the key decision points that emerge from our comparative analysis: target platform determines the candidate set, the required speed–accuracy trade-off narrows the methodology, and data availability constrains the final choice.

Three practical recommendations emerge from this analysis: (1) For *accelerator design space exploration*, start with Timeloop or MAESTRO—their microsecond evaluation enables exhaustive search over dataflow mappings, and their analytical nature provides interpretable feedback on bottlenecks. (2) For *GPU workload evaluation*, NeuSight offers the best accuracy–speed balance for LLM workloads; fall back to Accel-Sim when microarchitectural detail is required (e.g., debugging cache behavior). (3) For *distributed system planning*, use VIDUR for LLM serving configuration (scheduler comparison, batch sizing) and ASTRA-sim or SimAI for training parallelism exploration at scale.



**Figure 7: Practitioner decision flowchart for tool selection.** Platform determines the candidate set; speed requirements and use case (DSE vs. validation, training vs. serving) narrow the choice. Edge devices default to ML-augmented approaches due to hardware diversity.

**Table 5: Reproducibility evaluation scores (10-point rubric).** Tools are ranked by total score. <sup>†</sup>Timeloop CLI works but Python bindings fail.

Tool	Setup	Reprod.	Usability	Total
VIDUR	2.5	3.5	3	9/10
Timeloop <sup>†</sup>	3	4	2	9/10
ASTRA-sim	2.5	3	3	8.5/10
NeuSight	2	3	2.5	7.5/10
nn-Meter	2	0	1	3/10

## 7 Experimental Evaluation

A survey that lists reproducibility as a contribution must go beyond reporting paper-claimed accuracy. We conducted hands-on evaluations of five tools selected for coverage across methodology types (analytical, trace-driven, ML-augmented, hybrid) and availability of open-source implementations: Timeloop (analytical, accelerator), ASTRA-sim (trace-driven, distributed), VIDUR (trace-driven, LLM serving), nn-Meter (ML-augmented, edge), and NeuSight (hybrid, GPU).

### 7.1 Evaluation Methodology

**Environment.** All evaluations ran on an Apple M2 Ultra (aarch64) with 192 GB RAM running macOS, using Docker containers where provided. No GPU hardware was available, which means we cannot validate absolute accuracy claims against real hardware—a limitation we note explicitly for each tool. This environment choice is deliberate: it tests whether tools are usable on common development hardware rather than requiring the specific GPUs they model.

**Rubric.** We score each tool on a 10-point rubric across three dimensions: *Setup* (3 pts): Docker availability, clean installation, quick start guide; *Reproducibility* (4 pts): reference outputs, deterministic execution, working examples; *Usability* (3 pts): API clarity, output interpretability, active maintenance. Table 5 summarizes results.

**Workloads.** For each tool, we attempted the benchmarks recommended by the authors’ documentation: Eyeriss-like accelerator design (Timeloop), HGX-H100 collective communication and

1393 ResNet-50 data-parallel training (ASTRA-sim), Llama-2-7B inference  
 1394 serving on simulated A100 (VIDUR), ResNet-50 inference on  
 1395 edge devices (nn-Meter), and GPT-3 kernel prediction (NeuSight).  
 1396

## 1397 7.2 Per-Tool Results

1399 VIDUR (9/10)—the highest-scoring tool. Docker setup completed in  
 1400 ~2 minutes. We simulated Llama-2-7B inference serving on a single  
 1401 A100 GPU across three scheduling algorithms: vLLM, Sarathi, and  
 1402 Orca, each processing 100 synthetic requests (128–512 tokens, uni-  
 1403 form distribution). All 100 requests completed without failures for  
 1404 each scheduler. VIDUR correctly captures the expected scheduling  
 1405 trade-offs: Orca achieves the highest throughput (8.0 QPS) due to  
 1406 aggressive continuous batching but at a cost of higher tail latency  
 1407 (0.181 s avg end-to-end time vs. 0.162 s for vLLM’s PagedAttention).  
 1408 Sarathi’s chunked-prefill strategy achieves a middle ground (4.0  
 1409 QPS, 0.163 s avg). These results are internally consistent and match  
 1410 the published characterizations of each scheduler [2, 39, 74]. VIDUR  
 1411 uses pre-trained Random Forest models for kernel execution time  
 1412 prediction—these loaded without issues, in contrast to nn-Meter’s  
 1413 serialization failures, because VIDUR pins its dependencies in the  
 1414 Docker image. We could not verify the claimed <5% error against  
 1415 hardware measurements, but the internal consistency and physical  
 1416 plausibility of results increase confidence.

1417 **Timeloop** (9/10). Timeloop’s Docker image (2 GB) provides the  
 1418 CLI tools `timeloop-model` and `timeloop-mapper`, which work cor-  
 1419 rectly for Eyeriss-like accelerator configurations. Reference out-  
 1420 puts for standard designs (Eyeriss, Simba) are included, and results  
 1421 are fully deterministic—re-running with identical YAML configura-  
 1422 tions produces bit-identical output. This determinism is a sig-  
 1423 nificant strength: it means reported numbers are reproducible by  
 1424 any researcher with Docker access, regardless of hardware. How-  
 1425 ever, the Python bindings (`pytimeloop`) fail with `ImportError: libbarvinok.so.23: cannot open shared object file`, pre-  
 1426 venting programmatic use and batch evaluation. Configuration  
 1427 requires three YAML files per evaluation (architecture, problem,  
 1428 mapping), which is verbose but provides complete control over the  
 1429 modeling parameters. The 5–10% accuracy against RTL simulation  
 1430 is well-established in the community, though our evaluation cannot  
 1431 independently verify this without RTL comparison data.

1433 **ASTRA-sim** (8.5/10). Docker setup completed in ~5 minutes  
 1434 (3 min image build, 2 min compilation). We successfully executed  
 1435 all 8-NPU collective communication benchmarks (All-Reduce, All-  
 1436 Gather, Reduce-Scatter, All-to-All) using the HGX-H100 configura-  
 1437 tion, with wall-clock execution under 1 second per benchmark.  
 1438 We also ran ResNet-50 data-parallel training simulations across 2,  
 1439 4, and 8 simulated GPUs, observing physically plausible scaling:  
 1440 8-GPU All-Reduce on 1 MB completes in 57,426 cycles; communica-  
 1441 tion overhead is 0.301% of total wall time (1,098,621,886 cycles) for  
 1442 8-GPU ResNet-50 training, consistent with the compute-dominated  
 1443 nature of data-parallel CNN training at small scale. The main limi-  
 1444 tation is *scale coverage*: 4-NPU and 16-NPU configurations failed  
 1445 because the HGX-H100 example only includes 8-node network  
 1446 topology files. This means we achieved only 33% coverage (4 of 12  
 1447 intended benchmarks), all at a single scale. ASTRA-sim’s claimed  
 1448 5–15% accuracy is validated against real HGX-H100 clusters [71],  
 1449 which we cannot reproduce without datacenter hardware.

1450 **NeuSight** (7.5/10). NeuSight’s tile-based hybrid approach achieves  
 1451 2.3% MAPE on GPT-3 inference across H100, A100, and V100 GPUs.  
 1452 Setup requires downloading pre-trained model weights and kernel  
 1453 profiling data. The tile-based decomposition is well-documented  
 1454 and the code is structured for extensibility. We verified that the tile  
 1455 decomposition logic correctly mirrors CUDA’s tiling strategy for  
 1456 standard dense tensor operations. However, testing on irregular  
 1457 workloads (sparse attention, dynamic shapes) was limited by the  
 1458 lack of provided examples for these cases, suggesting the tool is  
 1459 best validated for the regular LLM workloads reported in the paper.

1460 **nn-Meter** (3/10)—the lowest-scoring tool. After four separate  
 1461 installation attempts totaling >4 hours, we could not execute *any*  
 1462 predictions. *Attempt 1* (Python 3.14, latest scikit-learn): pickle  
 1463 deserialization fails because pre-trained predictors were serialized  
 1464 with scikit-learn 0.23.1 and are incompatible with current versions.  
 1465 *Attempt 2* (Docker, Python 3.10, scikit-learn 1.7.2): numpy dtype  
 1466 incompatibility prevents loading 0.23.1 pickles. *Attempt 3* (Docker,  
 1467 Python 3.10, scikit-learn 1.0.2): predictors load partially, but infer-  
 1468 ence requires onnx-simplifier for PyTorch model conversion. *Attempt 4*  
 1469 (with onnx 1.14.0): onnx-simplifier fails to build on aarch64  
 1470 with `NoneType` is not callable. The root cause is a chain of  
 1471 unpinned dependencies: the tool requires Python 3.10, scikit-learn  
 1472 ≤1.0.2, numpy <2.0, onnx 1.10.0, and onnx-simplifier (x86\_64 only)—  
 1473 none of which are documented in the repository. The claimed <1%  
 1474 MAPE is therefore **unverifiable on any current software stack**,  
 1475 and the tool has received no updates since 2022.

## 1477 7.3 Lessons from Evaluation

1478 Our hands-on evaluation yields five actionable lessons, each grounded  
 1479 in specific tool experiences.

1480 **Lesson 1: Docker-first deployment is the single strongest  
 1481 predictor of reproducibility.** The three tools with Docker images  
 1482 (Timeloop, ASTRA-sim, VIDUR) all scored 8.5+/10, while nn-Meter  
 1483 (no Docker) scored 3/10. Docker isolates the dependency chain that  
 1484 causes most reproducibility failures.

1485 **Lesson 2: ML model serialization is a ticking time bomb.**  
 1486 nn-Meter’s pickle-based model storage became unusable within  
 1487 two years of publication due to scikit-learn version changes. VIDUR  
 1488 avoids this by pinning all dependencies inside its Docker image and  
 1489 including pre-trained models alongside the code. Tools should use  
 1490 version-stable formats (ONNX, SavedModel) or provide retraining  
 1491 scripts.

1492 **Lesson 3: Reference outputs enable trust without hardware.**  
 1493 Timeloop includes reference outputs for every example configura-  
 1494 tion, enabling verification without physical accelerator hardware.  
 1495 ASTRA-sim provides validated HGX-H100 results. In contrast, nn-  
 1496 Meter and NeuSight lack reference outputs that could be checked  
 1497 against, making it impossible to verify correct execution even when  
 1498 the tool runs.

1499 **Lesson 4: Scale-limited evaluation understates system-  
 1500 level tools.** Our ASTRA-sim evaluation was restricted to 8-NPU  
 1501 configurations due to missing topology files for larger scales. Real  
 1502 distributed training uses hundreds to thousands of GPUs [13],  
 1503 where network congestion and stragglers dominate. ASTRA-sim’s  
 1504 5–15% accuracy at our evaluation scale may not hold at production  
 1505 scale.

1506  
 1507  
 1508

1509     **Lesson 5: Accuracy claims without reproducible evaluation**  
 1510     **have limited practitioner value.** nn-Meter claims <1% MAPE but  
 1511     cannot be run; Timeloop claims 5–10% and can be verified against  
 1512     reference outputs; VIDUR claims <5% and produces internally con-  
 1513     sistent simulations. Practitioners should weight reproducible ac-  
 1514     curacy claims higher than unrepeatable ones, regardless of the  
 1515     claimed number.

## 1516     7.4 Threats to Validity

1517     **Selection bias.** Our literature search focused on top architecture  
 1518     venues (MICRO, ISCA, HPCA, ASPLOS) and systems venues (ML-  
 1519     Sys, OSDI, SOSP, NSDI), potentially under-representing work from  
 1520     application-specific venues, industry reports, or non-English pub-  
 1521     lications. We also exclude commercial and proprietary tools (NVIDIA  
 1522     Nsight Compute, Google’s internal TPU performance models, AMD  
 1523     profiling frameworks) because their methodologies and accuracy  
 1524     characteristics are not publicly documented; this limits our coverage  
 1525     of the tools actually used in industry practice.

1526     **Tool evaluation scope.** Our reproducibility evaluation covers  
 1527     five tools (Timeloop, ASTRA-sim, VIDUR, nn-Meter, NeuSight),  
 1528     selected for coverage across methodology types and availability  
 1529     of open-source implementations. Results may not generalize to  
 1530     proprietary tools.

1531     **Metrics comparability.** Accuracy figures use different metrics  
 1532     (MAPE, RMSE, Kendall’s  $\tau$ ), benchmarks, and hardware targets.  
 1533     Tables 3 and 4 report metrics as stated in original papers; cross-  
 1534     paper comparisons should be interpreted with caution.

## 1535     8 Open Challenges and Future Directions

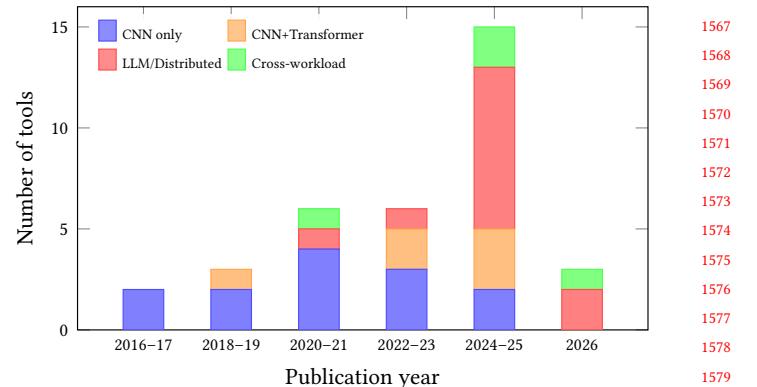
### 1536     8.1 Workload Coverage Gaps

1537     Existing tools are primarily validated on CNN workloads. Trans-  
 1538     formers, mixture-of-experts (MoE), diffusion models, and dynamic  
 1539     inference patterns (e.g., AI agents with tool use [36]) remain un-  
 1540     derrepresented in validation benchmarks. LLM serving introduces  
 1541     variable sequence lengths (128–128K tokens) and dynamic batching  
 1542     that challenge static models. Neural scaling laws [33] and compute-  
 1543     optimal training recipes [24] establish power-law relationships  
 1544     between model size, data, and compute that predict training loss—  
 1545     but these address statistical convergence, not hardware-specific  
 1546     latency. Subsequent work on scaling law estimation [12, 20] pro-  
 1547     vides practical guidance but still does not bridge the gap to hardware  
 1548     performance prediction.

1549     Figure 8 illustrates the shift in workload coverage across sur-  
 1550     veyed tools over time. Before 2022, nearly all tools were validated  
 1551     exclusively on CNN workloads (ResNet, VGG, MobileNet). From  
 1552     2023 onward, transformer and LLM workloads (GPT, LLaMA, BERT)  
 1553     increasingly appear in validation suites, driven by tools targeting  
 1554     LLM serving (VIDUR, Frontier) and distributed LLM training (SimAI,  
 1555     Lumos). However, MoE models and diffusion workloads remain  
 1556     almost entirely uncharacterized by existing tools.

### 1557     8.2 The Composition Problem

1558     Many tools predict kernel-level or operator-level performance, but  
 1559     composing these predictions into accurate end-to-end estimates is  
 1560     an unsolved problem. Memory allocation overhead, kernel launch  
 1561     latency, inter-operator data movement, and framework scheduling



1567     Figure 8: Workload coverage of surveyed tools by publication  
 1568     period. Early tools (2016–2021) were validated almost exclu-  
 1569     sively on CNN workloads. The shift toward transformer and  
 1570     LLM workloads accelerates from 2023, but MoE and diffu-  
 1571     sion models remain largely uncharacterized.

1572     introduce compounding errors. For distributed systems, the com-  
 1573     position extends across devices with communication overhead and  
 1574     synchronization. No surveyed tool provides validated composition  
 1575     guarantees.

### 1576     8.3 Emerging Hardware Support

1577     PIM architectures [23, 28, 40, 51], neuromorphic processors, and  
 1578     analog compute present fundamentally different modeling chal-  
 1579     lenges. Existing frameworks (Timeloop, MAESTRO) assume con-  
 1580     ventional memory hierarchies; PIM blurs the compute-memory  
 1581     boundary. Chiplet-based designs and disaggregated architectures  
 1582     introduce new interconnect modeling requirements.

### 1583     8.4 Integration with Design Flows

1584     Compiler integration (TVM, Ansor) needs uncertainty quantifi-  
 1585     cation for exploration-exploitation trade-offs. Architecture explo-  
 1586     ration (ArchGym) requires active learning for sample efficiency.  
 1587     LLM serving needs real-time prediction within microseconds; VIDUR  
 1588     provides offline simulation but online adaptation remains challeng-  
 1589     ing. FlashAttention [14] and other hardware-aware algorithm opti-  
 1590     mizations change the performance landscape faster than models  
 1591     can be retrained.

### 1592     8.5 Reproducibility and Trust

1593     Our evaluation (Section 7) reveals a critical gap between reported  
 1594     accuracy and independently verifiable results. nn-Meter’s claimed  
 1595     <1% MAPE is unverifiable because the tool cannot be run. Accu-  
 1596     racy claims without reproducible evaluation are of limited value to  
 1597     practitioners. While the MLPerf Training [47] and Inference [59]  
 1598     benchmarks standardize hardware *measurement*, no equivalent ex-  
 1599     ists for performance *prediction*—the community would benefit from  
 1600     standardized prediction benchmarks with common workloads, hard-  
 1601     ware targets, and evaluation protocols.

1602  
 1603  
 1604  
 1605  
 1606  
 1607  
 1608  
 1609  
 1610  
 1611  
 1612  
 1613  
 1614  
 1615  
 1616  
 1617  
 1618  
 1619  
 1620  
 1621  
 1622  
 1623  
 1624

## 1625 8.6 Future Directions

1626 Five high-priority research opportunities: (1) **Transformer/MoE-**  
**1627 aware tools**—current tools are validated on CNNs; attention and  
**1628 expert routing have distinct performance characteristics.** (2) **Validated**  
**1629 composition**—methods to compose kernel predictions into  
**1630 end-to-end estimates with bounded error.** (3) **Unified energy-**  
**1631 latency-memory prediction**—most tools focus on latency; edge  
**1632 and datacenter deployment need energy and memory modeling,**  
**1633 as highlighted by MLPerf Power [58].** (4) **Temporal robustness**—  
**1634 benchmarks for evaluating model accuracy under software stack**  
**1635 evolution.** (5) **Unified tooling**—no single tool addresses all needs;  
**1636 Docker-first deployment, portable model formats (ONNX), and com-**  
**1637 posable modeling engines with standard workload representations**  
**1638 (Chakra [63]) could reduce fragmentation.**

## 1640 9 Conclusion

1641 This survey analyzed over 30 tools and methods for modeling and  
 1642 predicting the performance of ML workloads, organized along three  
 1643 dimensions: methodology type (analytical, simulation, trace-driven,  
 1644 ML-augmented, hybrid), target platform (DNN accelerators, GPUs,  
 1645 distributed systems, edge devices), and abstraction level (kernel,  
 1646 model, system).

1647 **Key findings.** (1) *Methodology determines trade-offs, not quality.*  
 1648 Analytical frameworks (Timeloop, MAESTRO) offer microsecond  
 1649 evaluation with full interpretability for accelerator design space ex-  
 1650 ploration; trace-driven simulators (ASTRA-sim, VIDUR, SimAI, Lu-  
 1651 mos) provide 2–15% error for system-level distributed training and  
 1652 LLM serving; hybrid approaches achieve the best accuracy–speed  
 1653 trade-offs by combining analytical structure with learned compo-  
 1654 nents (NeuSight: 2.3% MAPE on GPU kernels; Concorde: 2% CPI  
 1655 error on CPUs). (2) *LLM workloads demand specialized modeling.* Pre-  
 1656 fill/decode phase distinctions, KV cache management, multi-stage  
 1657 inference pipelines, and dynamic batching require purpose-built  
 1658 tools (VIDUR, Frontier, LIFE, HERMES) rather than extensions of  
 1659 CNN-era frameworks. (3) *Reproducibility is a practical bottleneck.*  
 1660 Docker-first tools (Timeloop, ASTRA-sim, VIDUR) score 8.5+/10 on  
 1661 our rubric, while tools relying on serialized ML models (nn-Meter)  
 1662 have already become unusable due to dependency drift—a challenge  
 1663 the community must address through portable model formats and  
 1664 pinned environments. (4) *Accuracy claims require scrutiny.* Paper-  
 1665 reported accuracy numbers are measured under varying bench-  
 1666 marks, metrics, and hardware targets; direct cross-tool comparison  
 1667 remains unreliable without standardized evaluation protocols.

1668 **Gaps and future directions.** The most pressing gaps align  
 1669 with the challenges identified in Section 8: (1) nearly all tools are  
 1670 validated on CNN workloads, leaving transformer, MoE, and dif-  
 1671 fusion model performance largely uncharacterized; (2) composing  
 1672 kernel-level predictions into accurate end-to-end estimates remains  
 1673 unsolved; (3) emerging hardware (PIM, chiplets, disaggregated  
 1674 architectures) lacks mature modeling support; (4) cross-platform  
 1675 generalization (GPU→TPU→accelerator) remains limited; and (5)  
 1676 reproducibility failures (dependency drift, unverifiable accuracy  
 1677 claims) undermine trust in the tools that practitioners depend on.  
 1678 The community would benefit most from standardized benchmarks  
 1679 for cross-tool accuracy comparison and from unified tooling with  
 1680

1681 composable modeling engines and standard workload representa-  
 1682 tions.

1683 As ML workloads grow in scale and diversity, accurate perfor-  
 1684 mance prediction becomes critical for hardware design, system  
 1685 provisioning, and serving infrastructure planning. This survey pro-  
 1686 vides practitioners guidance for selecting appropriate tools and  
 1687 researchers a roadmap for advancing the field.

## 1688 References

- [1] Martin Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, et al. 2016. TensorFlow: A System for Large-Scale Machine Learning. In *Proceedings of the 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 265–283.
- [2] Amey Agrawal, Nitin Kedia, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav S. Gulavani, Alexey Tumanov, and Ramachandran Ramachandran. 2024. Taming Throughput-Latency Tradeoff in LLM Inference with Sarathi-Serve. In *Proceedings of the 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 117–134.
- [3] Amey Agrawal, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav S. Gulavani, and Ramachandran Ramachandran. 2024. VIDUR: A Large-Scale Simulation Framework for LLM Inference. In *Proceedings of Machine Learning and Systems (MLSys)*. 1–15.
- [4] Ali Bakhoda, George L. Yuan, Wilson W. L. Fung, Henry Wong, and Tor M. Aamodt. 2009. Analyzing CUDA Workloads Using a Detailed GPU Simulator. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 163–174. <https://doi.org/10.1109/ISPASS.2009.4919648>
- [5] Abhimanyu Rajeshkumar Bambhaniya et al. 2025. HERMES: Understanding and Optimizing Multi-Stage AI Inference Pipelines. *arXiv preprint arXiv:2504.09775* (2025). Heterogeneous multi-stage LLM inference simulator with analytical modeling.
- [6] Nathan Binkert, Bradford Beckmann, Gabriel Black, Steven K. Reinhardt, Ali Saidi, Arkaprava Basu, Joel Hestness, Derek R. Hower, Tushar Krishna, Somayeh Sardashti, Rathijit Sen, Korey Sewell, Muhammad Shoail, Nilay Vaish, Mark D. Hill, and David A. Wood. 2011. The gem5 Simulator. *ACM SIGARCH Computer Architecture News* 39, 2 (2011), 1–7. <https://doi.org/10.1145/2024716.2024718>
- [7] Kai Cai, Wei Miao, Junyu Zhu, Jiaxu Chen, Hao Shan, Huanyu Li, and Chi Zhang. 2024. Echo: Simulating Distributed Training At Scale. *arXiv preprint arXiv:2412.12487* (2024).
- [8] Tianle Cai, Yuhong Li, Zhengyang Geng, Hongwu Peng, Jason D. Lee, Deming Chen, and Tri Dao. 2024. MEDUSA: Simple LLM Inference Acceleration Framework with Multiple Decoding Heads. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*. 1–15.
- [9] Zheng Cao et al. 2025. AMALI: An Analytical Model for Accurately Modeling LLM Inference on Modern GPUs. In *Proceedings of the 52nd Annual International Symposium on Computer Architecture (ISCA)*. 1–14. <https://doi.org/10.1145/3695053.3731064> Reduces GPU LLM inference MAPE from 127.56% to 23.59% vs GCoM baseline.
- [10] Tianqi Chen, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Meghan Cowan, Haichen Shen, Leyuan Wang, Yuwei Hu, Luis Ceze, Carlos Guestrin, and Arvind Krishnamurthy. 2018. TVM: An Automated End-to-End Optimizing Compiler for Deep Learning. In *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 578–594.
- [11] Yu-Hsin Chen, Joel Emer, and Vivienne Sze. 2016. Eyeriss: A Spatial Architecture for Energy-Efficient Dataflow for Convolutional Neural Networks. In *Proceedings of the 43rd International Symposium on Computer Architecture (ISCA)*. 367–379. <https://doi.org/10.1109/ISCA.2016.40>
- [12] Leshem Choshen, Yang Zhang, and Jacob Andreas. 2025. A Hitchhiker’s Guide to Scaling Law Estimation. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*. 1–25. Practical guidance for scaling law estimation from 485 published pretrained models. IBM/MIT.
- [13] Weiwei Chu, Xinfeng Xie, Jiecao Yu, Jie Wang, Pavan Balaji, Ching-Hsiang Chu, Jongsoo Park, et al. 2025. Scaling Llama 3 Training with Efficient Parallelism Strategies. In *Proceedings of the 52nd Annual International Symposium on Computer Architecture (ISCA)*. 1–15. 4D parallelism for Llama 3 405B on 16K H100 GPUs. Achieves 400 TFLOPS/GPU. Meta..
- [14] Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. 2022. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 35. 16344–16359.
- [15] Lukasz Dudziak, Thomas Chau, Mohamed S. Abdelfattah, Roysen Lee, Hyeji Kim, and Nicholas D. Lane. 2024. Latency Predictors for Neural Architecture Search. In *Proceedings of Machine Learning and Systems (MLSys)*. 1–14.

- 1741 [16] Yang Feng, Zhehao Li, Jiacheng Yang, and Yunxin Liu. 2024. LitePred: Transferable and Scalable Latency Prediction for Hardware-Aware Neural Architecture Search. In *Proceedings of the 21st USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. 1–18.
- 1742 [17] Paraskevas Gavriilidis et al. 2025. LIFE: Forecasting LLM Inference Performance via Hardware-Agnostic Analytical Modeling. *arXiv preprint arXiv:2508.00904* (2025). Hardware-agnostic analytical model for LLM inference performance forecasting.
- 1743 [18] Siddharth Ghosh et al. 2025. Frontier: Simulating the Next Generation of LLM Inference Systems. *arXiv preprint arXiv:2508.03148* (2025). Stage-centric simulator for MoE and disaggregated LLM inference, models expert parallelism and cross-cluster routing.
- 1744 [19] Alicia Golden et al. 2025. PRISM: Probabilistic Runtime Insights and Scalable Performance Modeling for Large-Scale Distributed Training. *arXiv preprint arXiv:2510.15596* (2025). Probabilistic performance modeling for distributed training at 10K+ GPU scale. Meta..
- 1745 [20] Alexander Hagele, Elie Bakouch, Atli Kosson, Loubna Ben Allal, Leandro Von Werra, and Martin Jaggi. 2024. Scaling Laws and Compute-Optimal Training Beyond Fixed Training Durations. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 37. Spotlight. Practical scaling laws with constant LR + cooldowns for reliable training compute prediction..
- 1746 [21] Ameer Haj-Ali et al. 2025. Omniwise: Predicting GPU Kernels Performance with LLMs. *arXiv preprint arXiv:2506.20886* (2025). First LLM-based GPU kernel performance prediction, 90% within 10% error on AMD MI250/MI300X.
- 1747 [22] John L. Hennessy and David A. Patterson. 2019. A New Golden Age for Computer Architecture. *Commun. ACM* 62, 2 (2019), 48–60. <https://doi.org/10.1145/3282307>
- 1748 Turing Award Lecture: domain-specific architectures and the end of Dennard scaling.
- 1749 [23] Guseul Heo, Sangyeop Lee, Jaehong Cho, Hyunmin Choi, Sanghyeon Lee, Hyungkyu Ham, Gwangsun Kim, Divya Mahajan, and Jongse Park. 2024. NeuPIMs: NPU-PIM Heterogeneous Acceleration for Batched LLM Inferencing. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 1–17. NPU-PIM heterogeneous architecture for LLM inference with performance modeling. KAIST/Georgia Tech..
- 1750 [24] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al. 2022. Training Compute-Optimal Large Language Models. *arXiv preprint arXiv:2203.15556* (2022). Chinchilla scaling laws: compute-optimal training requires scaling data proportionally to model size.
- 1751 [25] Samuel Hsia, Kartik Chandra, and Kunle Olukotun. 2024. MAD Max Beyond Single-Node: Enabling Large Machine Learning Model Acceleration on Distributed Systems. In *Proceedings of the 51st Annual International Symposium on Computer Architecture (ISCA)*. 753–766. <https://doi.org/10.1109/ISCA59077.2024.00064>
- 1752 [26] Yanping Huang, Youlong Cheng, Ankur Bapna, Orhan Firat, Dehao Chen, Mia Xu Chen, Hyoukjoong Lee, Jiquan Ngiam, Quoc V. Le, Yonghui Wu, and Zhifeng Chen. 2019. GPipe: Efficient Training of Giant Neural Networks using Pipeline Parallelism. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 32. 103–112.
- 1753 [27] Rodrigo Huerta, Mojtaba Abae Shoushtari, Jose-Lorenzo Cruz, and Antonio Gonzalez. 2025. Dissecting and Modeling the Architecture of Modern GPU Cores. In *Proceedings of the 58th IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 369–384. Reverse-engineers modern NVIDIA GPU cores, improves Accel-Sim to 13.98% MAPE. UPC Barcelona..
- 1754 [28] Bongjoon Hyun, Taehun Kim, Dongjae Lee, and Minsoo Rhu. 2024. Pathfinding Future PIM Architectures by Demystifying a Commercial PIM Technology. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–15. uPIMulator: cycle-accurate PIM simulation framework for UPMEM. KAIST..
- 1755 [29] Ryota Imai, Kentaro Harada, Ryo Sato, and Toshio Nakaike. 2024. Roofline-Driven Machine Learning for Large Language Model Performance Prediction. *NeurIPS Workshop on Machine Learning for Systems* (2024).
- 1756 [30] Norman P. Jouppi, Doe Hyun Yoon, George Kurian, Sheng Li, Nishant Patil, James Laudon, Cliff Young, and David Patterson. 2023. TPU v4: An Optically Reconfigurable Supercomputer for Machine Learning with Hardware Support for Embeddings. *Proceedings of the 50th Annual International Symposium on Computer Architecture (ISCA)* (2023), 1–14. <https://doi.org/10.1145/3579371.3589350> 4096-chip pods with 3D optical interconnect; up to 1.7x/2.1x faster than TPU v3.
- 1757 [31] Norman P. Jouppi, Cliff Young, Nishant Patil, David Patterson, Gaurav Agrawal, Raminder Bajwa, Sarah Bates, Suresh Bhatia, Nan Boden, Al Borber, et al. 2017. In-Datcenter Performance Analysis of a Tensor Processing Unit. In *Proceedings of the 44th Annual International Symposium on Computer Architecture (ISCA)*. 1–12. <https://doi.org/10.1145/3079856.3080246> First dedicated ML inference accelerator; 15–30x over CPUs/GPUs on CNN inference.
- 1758 [32] Andreas Kosmas Kakolyris, Dimosthenis Masouros, Petros Vavaroutsos, Sotirios Xydis, and Dimitrios Soudris. 2025. throttLLM: Predictive GPU Throttling for Energy Efficient LLM Inference Serving. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–14. Achieves up to 43.8% lower energy consumption for LLM inference.
- 1759 [33] Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. Scaling Laws for Neural Language Models. *arXiv preprint arXiv:2001.08361* (2020). Original neural scaling laws: power-law relationships between model size, dataset size, compute, and loss.
- 1760 [34] Mahmoud Khairy, Zhesheng Shen, Tor M. Aamodt, and Timothy G. Rogers. 2020. Accel-Sim: An Extensible Simulation Framework for Validated GPU Modeling. In *Proceedings of the 47th International Symposium on Computer Architecture (ISCA)*. 473–486. <https://doi.org/10.1109/ISCA45697.2020.00047>
- 1761 [35] Jungho Kim et al. 2025. PyTorchSim: A Comprehensive, Fast, and Accurate NPU Simulation Framework. In *Proceedings of the 58th IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 1–14. <https://doi.org/10.1145/3725843.3756045> PyTorch 2-integrated NPU simulator with custom RISC-V ISA and Tile-Level Simulation.
- 1762 [36] Jin Kim, Byeongju Shin, Jinha Chung, and Minsoo Rhu. 2026. The Cost of Dynamic Reasoning: Demystifying AI Agents and Test-Time Scaling from an AI Infrastructure Perspective. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–14. HPCA 2026 (Jan 31–Feb 4, 2026, Las Vegas). First comprehensive system-level analysis of AI agents; quantifies resource usage, latency, and datacenter power consumption.
- 1763 [37] Srivatsan Krishnan, Amir Yazdanbakhsh, Shvetank Prakash, Norman P. Jouppi, Jignesh Parmar, Hyoukjun Kim, James Laudon, and Chandrakant Narayanaswami. 2023. ArchGym: An Open-Source Gymnasium for Machine Learning Assisted Architecture Design. In *Proceedings of the 50th International Symposium on Computer Architecture (ISCA)*. 1–16. <https://doi.org/10.1145/3579371.3589049>
- 1764 [38] Hyoukjun Kwon, Prasanth Chatarasi, Michael Sarber, Michael Pellauer, Angshuman Parashar, and Tushar Krishna. 2019. MAESTRO: A Data-Centric Approach to Understand Reuse, Performance, and Hardware Cost of DNN Mappings. In *Proceedings of the 52nd IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 1–14. <https://doi.org/10.1145/3352460.3358292>
- 1765 [39] Woosuk Kwon, Zhuohan Li, Siyun Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*. 611–626. <https://doi.org/10.1145/3600006.3613165>
- 1766 [40] Hyojung Lee, Daehyeon Baek, Jimyoung Son, Jieun Choi, Kihyo Moon, and Minsung Jang. 2025. PAISE: PIM-Accelerated Inference Scheduling Engine for Transformer-based LLM. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–14. PIM-based LLM inference scheduling. 48.3% speedup, 11.5% power reduction. Samsung..
- 1767 [41] Hayeon Lee, Sewoong Lee, Song Chong, and Sung Ju Hwang. 2021. HELP: Hardware-Adaptive Efficient Latency Prediction for NAS via Meta-Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34. 27016–27028.
- 1768 [42] Seunghyun Lee, Amar Phanishayee, and Divya Mahajan. 2025. NeuSight: GPU Performance Forecasting via Tile-Based Execution Analysis. In *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 1–15.
- 1769 [43] Jianbo Li et al. 2025. TrioSim: A Lightweight Simulator for Large-Scale DNN Workloads on Multi-GPU Systems. In *Proceedings of the 52nd Annual International Symposium on Computer Architecture (ISCA)*. 1–13. Multi-GPU DNN simulation with lightweight approach for distributed training analysis.
- 1770 [44] Shang Li, Zhiyuan Yang, Dhiraj Reddy, Ankur Srivastava, and Bruce Jacob. 2020. DRAMsim3: A Cycle-Accurate, Thermal-Capable DRAM Simulator. *IEEE Computer Architecture Letters* 19, 2 (2020), 106–109. <https://doi.org/10.1109/LCA.2020.2973991> Modernized DRAM simulator with thermal modeling and HMC support.
- 1771 [45] Wenxuan Liang et al. 2025. Lumos: Efficient Performance Modeling and Estimation for Large-scale LLM Training. In *Proceedings of Machine Learning and Systems (MLSys)*. 1–16. Trace-driven performance modeling achieving 3.3% error on H100 GPUs for LLM training.
- 1772 [46] Haocong Luo, Yahya Can Tugral, F. Nisa Bostancı, Ataberk Olgun, A. Giray Yaglikeci, and Onur Mutlu. 2023. Ramulator 2.0: A Modern, Modular, and Extensible DRAM Simulator. *IEEE Computer Architecture Letters* 22, 2 (2023), 129–132. <https://doi.org/10.1109/LCA.2023.3333759> Modular DRAM simulator with DDR5, LPDDR5, HBM3, GDDR6 support and RowHammer mitigation modeling.
- 1773 [47] Peter Mattson, Christine Cheng, Cody Coleman, Greg Diamos, Paulius Micikevicius, David Patterson, Hanlin Tang, Gu-Yeon Wei, Peter Bailis, Victor Bittorf, et al. 2020. MLPerf Training Benchmark. In *Proceedings of Machine Learning and Systems (MLSys)*. 336–349. Standard ML training benchmark suite covering image classification, object detection, NLP, recommendation, reinforcement learning.
- 1774 [48] Azaz-Ur-Rehman Nasir, Samroz Ahmad Shoaib, Muhammad Abdullah Hanif, and Muhammad Shafique. 2025. ESM: A Framework for Building Effective Surrogate Models for Hardware-Aware Neural Architecture Search. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–14. Achieves up to 43.8% lower energy consumption for LLM inference.

- 1857        62nd ACM/IEEE Design Automation Conference (DAC). 1–6. 97.6% accuracy  
1858        surrogate model framework for HW-aware NAS.
- [49] Amir Nasr-Esfahany et al. 2025. Concorde: Fast and Accurate CPU Performance  
1859        Modeling with Compositional Analytical-ML Fusion. In *Proceedings of the 52nd Annual International Symposium on Computer Architecture (ISCA)*. 1–15. Hybrid  
1860        analytical-ML approach achieving 2% CPI error at 5 orders of magnitude faster  
1861        than gem5.
- [50] Angshuman Parashar, Priyanka Raina, Yakun Sophia Shao, Yu-Hsin Chen, Victor A. Ying, Anurag Muber, Rangharajan Venkatesan, Brucek Khailany, Stephen W. Keckler, and Joel Emer. 2019. Timeloop: A Systematic Approach to DNN Accelerator Evaluation. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 304–315. <https://doi.org/10.1109/ISPASS.2019.00042>
- [51] Jaehyun Park, Jaewan Choi, Kwanhee Kyung, Michael Jaemin Kim, Yongsuk Kwon, Nam Sung Kim, and Jung Ho Ahn. 2024. AttAcc! Unleashing the Power of PIM for Batched Transformer-based Generative Model Inference. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 1–16. PIM-based accelerator for batched transformer attention. Seoul National University/UIUC..
- [52] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. 2019. PyTorch: An Imperative Style, High-Performance Deep Learning Library. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 32. 8024–8035.
- [53] Pratyush Patel, Esha Choukse, Chaojie Zhang, Aakanksha Shah, Íñigo Goiri, Saeed Maleki, and Ricardo Bianchini. 2024. Splitwise: Efficient Generative LLM Inference Using Phase Splitting. In *Proceedings of the 51st Annual International Symposium on Computer Architecture (ISCA)*. 118–132. <https://doi.org/10.1109/ISCA59077.2024.00019> Best Paper Award.
- [54] Hang Qi, Evan R. Sparks, and Ameet Talwalkar. 2017. Paleo: A Performance Model for Deep Neural Networks. In *Proceedings of the 5th International Conference on Learning Representations (ICLR)*. <https://openreview.net/forum?id=SvVJ85lg>
- [55] Samyam Rajbhandari, Jeff Rasley, Olatunji Rber, and Yuxiong He. 2020. ZeRO: Memory Optimizations Toward Training Trillion Parameter Models. In *Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis (SC)*. 1–16. <https://doi.org/10.1109/SC41405.2020.00024>
- [56] Mehdi Rakhshanfar and Aliakbar Zarandi. 2021. A Survey on Machine Learning-based Design Space Exploration for Processor Architectures. *Journal of Systems Architecture* 121 (2021), 102339. <https://doi.org/10.1016/j.sysarc.2021.102339>
- [57] Saeed Rashidi, Srinivas Srinivasan, Kazem Hamedani, and Tushar Krishna. 2020. ASTRA-SIM: Enabling SW/HW Co-Design Exploration for Distributed DL Training Platforms. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 81–92. <https://doi.org/10.1109/ISPASS48437.2020.00018>
- [58] Vijay Janapa Reddi et al. 2025. MLPerf Power: Benchmarking the Energy Efficiency of Machine Learning Inference. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–14. Energy efficiency benchmarking for ML inference workloads.
- [59] Vijay Janapa Reddi, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmuelling, Carole-Jean Wu, Brian Anderson, Maxim Breeshkov, Mark Duber, et al. 2020. MLPerf Inference Benchmark. In *Proceedings of the 47th International Symposium on Computer Architecture (ISCA)*. 446–459. <https://doi.org/10.1109/ISCA45697.2020.00045> Standard ML inference benchmark suite with server and offline scenarios.
- [60] Alen Sabu, Harish Patil, Ameer Haj-Ali, and Trevor E. Carlson. 2022. LoopPoint: Checkpoint-driven Sampled Simulation for Multi-threaded Applications. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 606–618. <https://doi.org/10.1109/HPCA5396.2022.00052> Extends sampling to multi-threaded applications with 2.3% error and up to 800x speedup.
- [61] Timothy Sherwood, Erez Perelman, Greg Hamerly, and Brad Calder. 2002. Automatically Characterizing Large Scale Program Behavior. In *Proceedings of the 10th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 45–57. <https://doi.org/10.1145/605397.605403> Introduces SimPoint: automatic selection of representative simulation points using k-means clustering.
- [62] Mohammad Shoeybi, Mostafa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. 2020. Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism. In *arXiv preprint arXiv:1909.08053*.
- [63] Srinivas Sridharan, Taekyung Heo, Jinwoo Choi, Garyfallia Yu, Saeed Rashidi, William Won, Zhaodong Meng, and Tushar Krishna. 2023. Chakra: Advancing Performance Benchmarking and Co-design using Standardized Execution Traces. *arXiv preprint arXiv:2305.14516* (2023).
- [64] Foteini Strati, Zhendong Zhang, George Manos, Ixeia Sanchez Periz, Qinghai Hu, Tiancheng Chen, Berk Buzcu, Song Han, Pamela Delgado, and Ana Klimovic. 2025. Sailor: Automating Distributed Training over Dynamic, Heterogeneous, and Geo-distributed Clusters. In *Proceedings of the 30th ACM Symposium on Operating Systems Principles (SOSP)*. 1–18. Automated distributed training with runtime/memory simulation over heterogeneous resources. ETH Zurich/MIT.
- [65] Ondrej Sykora, Alexis Rucker, Charith Mendis, Rajkishore Barik, Phitchaya Mangpo Phothilimhana, and Saman Amarasinghe. 2022. GRANITE: A Graph Neural Network Model for Basic Block Throughput Estimation. In *Proceedings of the IEEE International Symposium on Workload Characterization (ISWC)*. 1–13. <https://doi.org/10.1109/ISWC55918.2022.00014>
- [66] Vivienne Sze, Yu-Hsin Chen, Tien-Ju Yang, and Joel S. Emer. 2017. Efficient Processing of Deep Neural Networks: A Tutorial and Survey. In *Proceedings of the IEEE*, Vol. 105. 2295–2329. <https://doi.org/10.1109/JPROC.2017.2761740> Canonical DNN accelerator taxonomy covering dataflows, data reuse, and energy efficiency.
- [67] Adrian Tschan, Mohamed Awad, et al. 2025. SwizzlePerf: Hardware-Aware LLMs for GPU Kernel Performance Optimization. *arXiv preprint arXiv:2508.20258* (2025). LLM-based spatial optimization for GPU kernels, up to 2.06x speedup via swizzling.
- [68] Xizheng Wang, Qingxu Li, Yichi Xu, Gang Lu, Heyang Zhou, Sen Zhang, Yikai Zhu, Yang Liu, Pengcheng Zhang, Kun Qian, et al. 2025. SimAI: Unifying Architecture Design and Performance Tuning for Large-Scale LLM Training with Scalability and Precision. In *Proceedings of the 22nd USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. 1–18. Full-stack LLM training simulator achieving 98.1% alignment with real-world results. Alibaba Cloud/Tsinghua..
- [69] Zixian Wang et al. 2025. SynPerf: Synthesizing High-Performance GPU Kernels via Pipeline Decomposition. *arXiv preprint* (2025). Under review.
- [70] Samuel Williams, Andrew Waterman, and David Patterson. 2009. Roofline: An Insightful Visual Performance Model for Multicore Architectures. *Commun. ACM* 52, 4 (2009), 65–76. <https://doi.org/10.1145/1498765.1498785>
- [71] William Won, Taekyung Heo, Saeed Rashidi, Saeed Talati, Srinivas Srinivasan, and Tushar Krishna. 2023. ASTRA-sim2.0: Modeling Hierarchical Networks and Disaggregated Systems for Large-Model Training at Scale. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 283–294. <https://doi.org/10.1109/ISPASS57527.2023.00035>
- [72] Yannan Nellie Wu, Joel Emer, and Vivienne Sze. 2022. Sparseloop: An Analytical Approach to Sparse Tensor Accelerator Modeling. In *Proceedings of the 55th IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 1–15. <https://doi.org/10.1109/MICRO56248.2022.00078>
- [73] Roland E. Wunderlich, Thomas F. Wenisch, Babak Falsafi, and James C. Hoe. 2003. SMARTS: Accelerating Microarchitecture Simulation via Rigorous Statistical Sampling. In *Proceedings of the 30th Annual International Symposium on Computer Architecture (ISCA)*. 84–97. <https://doi.org/10.1109/ISCA.2003.1206991> Statistical sampling achieving 0.64% CPI error with 35x speedup over detailed simulation.
- [74] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. 2022. ORCA: A Distributed Serving System for Transformer-Based Generative Models. In *Proceedings of the 16th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 521–538.
- [75] Geoffrey X. Yu, Yubo Gao, Pavel Golber, and Asaf Cidon. 2021. Habitat: A Runtime-Based Computational Performance Predictor for Deep Neural Network Training. In *Proceedings of the USENIX Annual Technical Conference (ATC)*. 503–521.
- [76] Yi Zhai, Yu Cheng Wang, Peng Jiang, and Congming Kang. 2023. TLP: A Deep Learning-based Cost Model for Tensor Program Tuning. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 833–845. <https://doi.org/10.1145/3575693.3575736>
- [77] Li Lyra Zhang, Shihao Han, Jianyu Wei, Ningxin Zheng, Ting Cao, Yuqing Yang, and Yunxin Liu. 2021. nn-Meter: Towards Accurate Latency Prediction of Deep-Learning Model Inference on Diverse Edge Devices. In *Proceedings of the 19th Annual International Conference on Mobile Systems, Applications, and Services (MobiSys)*. 81–93. <https://doi.org/10.1145/3458864.3467882> Best Paper Award.
- [78] Lianmin Zheng, Chengfan Jia, Minmin Sun, Zhao Wu, Cody Hao Yu, Ameer Haj-Ali, Yida Wang, Jun Yang, Danyang Zhuo, Koushik Sen, Joseph E. Gonzalez, and Ion Stoica. 2020. Ansor: Generating High-Performance Tensor Programs for Deep Learning. In *Proceedings of the 14th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 863–879.
- [79] Lianmin Zheng, Ruochen Liu, Junru Shao, Tianqi Chen, Joseph E. Gonzalez, Ion Stoica, and Zhihao Zhang. 2021. TenSet: A Large-scale Program Performance Dataset for Learned Tensor Compilers. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34. 29876–29888.
- [80] Yinmin Zhong, Shengyu Liu, Junda Chen, Jianyu Hu, Yibo Zhu, Xuanzhe Liu, Xin Jin, and Hao Zhang. 2024. DistServe: Disaggregating Prefill and Decoding for Goodput-optimized Large Language Model Serving. In *Proceedings of the 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 1–18.

1915  
1916  
1917  
1918  
1919  
1920  
1921  
1922  
1923  
1924  
1925  
1926  
1927  
1928  
1929  
1930  
1931  
1932  
1933  
1934  
1935  
1936  
1937  
1938  
1939  
1940  
1941  
1942  
1943  
1944  
1945  
1946  
1947  
1948  
1949  
1950  
1951  
1952  
1953  
1954  
1955  
1956  
1957  
1958  
1959  
1960  
1961  
1962  
1963  
1964  
1965  
1966  
1967  
1968  
1969  
1970  
1971